

ets_marl - Default RQ Analysis

May 8, 2026

1 ETS MARL — Default-Config Research-Question Analysis

This notebook focuses on what the **default configuration** alone can tell us about the thesis research questions. It deliberately does **not** attempt cross-config (variant) comparisons — those will be addressed in a separate experiments notebook.

Research questions (full text in docs/research_questions.md and the top-level README.md):

- **Main RQ** — How can MARL be used to simulate the EU ETS, and what does the simulation reveal about its regulatory mechanisms and the strategic behaviour of power-sector participants?
- **Sub-RQ 1** — What environment / algorithm design choices are required to build a *stable, behaviourally credible* simulation?
- **Sub-RQ 2** — How sensitive are market outcomes (price stability, compliance, decarbonisation pace) to the regulatory mechanisms of the EU ETS? (*Default-only descriptive baseline; cross-variant claims deferred.*)
- **Sub-RQ 3** — How do financial vs environmental objectives shape company strategies, and how do those aggregate into market-level outcomes?

Structure

1. Setup — locate project root.
2. Imports.
3. Multi-sweep loader (allows pulling in **several** default-config sweeps at once — extends the single-spec pattern from ets_marl - Sweep Analysis.ipynb).
4. Discovery & loading.
5. Sanity overview (counts, episode coverage, default-only filter).
6. **§RQ1** — convergence, cross-seed reproducibility, behavioural credibility, pathology.
7. **§RQ2** — within-default regulatory mechanism behaviour (cap → price, MSR / TNAC, penalty incidence, decarbonisation pace).
8. **§RQ3** — strategic behaviour by archetype and reward weighting, aggregated to system-level outcomes.

The notebook is intentionally compact: each section ships the headline plot or table that supports the corresponding RQ claim, and stops there.

```
[1]: # Install notebook dependencies if needed.
      # Set INSTALL_NOTEBOOK_DEPS = True (or run the cell with required packages
      ↪missing)
      # to pip-install everything in notebooks/requirements-notebooks.txt using the
```

```

# CURRENT kernel's interpreter. This is more reliable than `!pip install -r
↳ {path}`
# (IPython magics don't expand `{var}` at runtime) and works on Azure / Colab /
↳ local.
import subprocess, sys, importlib, re
from pathlib import Path

INSTALL_NOTEBOOK_DEPS = False # flip to True to force-reinstall

# Map PyPI distribution name -> import module name when they differ.
_DIST_TO_IMPORT = {
    "pyyaml": "yaml",
    "scikit-learn": "sklearn",
    "imbalanced-learn": "imblearn",
    "opencv-python": "cv2",
    "pillow": "PIL",
    "beautifulsoup4": "bs4",
    "ipython": "IPython",
}

def _find_requirements():
    req_name = "requirements-notebooks.txt"
    here = Path.cwd()
    candidates = [here / req_name, here / "notebooks" / req_name]
    candidates.extend(p / "notebooks" / req_name for p in here.parents)
    return next((p for p in candidates if p.exists()), None)

def _parse_requirements(req_path):
    """Extract distribution names from a requirements file (skip comments,
    ↳ options, env markers)."""
    dists = []
    for raw in req_path.read_text(encoding="utf-8").splitlines():
        line = raw.strip()
        if not line or line.startswith("#") or line.startswith("-"):
            continue
        # Strip env markers and extras: "pkg[extra]>=1.0; python_version>='3.9'"
        line = line.split(";", 1)[0].strip()
        m = re.match(r"^([A-Za-z0-9_\.\\-]+)", line)
        if m:
            dists.append(m.group(1).lower())
    return dists

def _missing_required(dists):
    missing = []
    for dist in dists:
        mod = _DIST_TO_IMPORT.get(dist, dist.replace("-", "_"))
        try:

```

```

        importlib.import_module(mod)
    except ImportError:
        missing.append(dist)
    return missing

req_path = _find_requirements()
if req_path is None:
    raise FileNotFoundError("Could not find notebooks/requirements-notebooks.
↳txt")

dists = _parse_requirements(req_path)
missing = _missing_required(dists)

if INSTALL_NOTEBOOK_DEPS or missing:
    if missing:
        print(f"Missing packages: {missing}")
        print(f"Installing notebook dependencies from: {req_path}")
        subprocess.check_call([sys.executable, "-m", "pip", "install", "-r",
↳str(req_path)])
    else:
        print(f"All {len(dists)} notebook dependencies present; skipping install.")
        print("(Set INSTALL_NOTEBOOK_DEPS = True to force a reinstall.)")

```

All 14 notebook dependencies present; skipping install.
(Set INSTALL_NOTEBOOK_DEPS = True to force a reinstall.)

1.1 1. Setup — Locate Project Root

```

[2]: import os, sys
from pathlib import Path

def _is_project_root(path: str) -> bool:
    return os.path.isdir(os.path.join(path, "src")) and os.path.isdir(os.path.
↳join(path, "configs"))

_candidates = ["..", ".", os.path.join("..", ".."),
               os.path.join("Thesis-Energy-Auction", "ets_marl_happo_current"),
↳"Thesis-Energy-Auction"]
_found = _is_project_root(".")
if not _found:
    for _c in _candidates:
        if _is_project_root(_c):
            os.chdir(_c); _found = True; break
if not _found:
    raise FileNotFoundError(f"Cannot find project root (needs src/ + configs/).
↳\nCWD: {os.getcwd()} | Contents: {os.listdir('.')}")

```

```
PROJECT_ROOT = Path.cwd().resolve()
sys.path.insert(0, str(PROJECT_ROOT))
print(f"Working directory: {PROJECT_ROOT}")
```

Working directory: C:\Users\danie\Documents\Python_Scripts\Master Thesis\Thesis-Energy-Auction

1.2 2. Imports

```
[3]: import re
import yaml
import warnings
from pathlib import Path

import numpy as np
import pandas as pd
from src.utils.run_data import load_run_csv, glob_run_logs # parquet-cached
↳ log loader
import matplotlib.pyplot as plt
import seaborn as sns

sns.set_style('whitegrid')
pd.set_option('display.max_columns', 200)
pd.set_option('display.width', 200)

from src.utils.sweep import load_sweep_spec, expand_dotted_overrides, deep_merge
```

1.3 3. Sweep Specs to Load

List one or more sweep specs (paths relative to the project root). All of them are loaded and merged into a single dict keyed by (`sweep_name`, `variant_name`, `seed`). The default below loads only `default_seeds.yaml`, but additional default-config sweeps can be appended freely — every run that survives the §5 filter is treated as another seed of the *same* default-config population.

If a sweep contains variants whose `overrides` deviate from the base config, those variants are flagged in §5 and excluded from the analysis by default (this notebook is for the default config only). To include them anyway, set `STRICT_DEFAULT_ONLY = False` in §5.

```
[4]: # Add additional default-config sweeps here as the project produces them.
SWEEP_SPECS = [
    PROJECT_ROOT / 'configs' / 'sweeps' / 'default_seeds.yaml',
    # PROJECT_ROOT / 'configs' / 'sweeps' / 'another_default_sweep.yaml',
]

# Optional: override the output_dir from the spec (useful if results were
# moved). One entry per spec, or None to use the spec's own output_dir.
OUTPUT_DIR_OVERRIDES = []
```

```

PROJECT_ROOT / 'results' / 'sweeps' / 'default',
# None,
]
assert len(OUTPUT_DIR_OVERRIDES) == len(SWEEP_SPECS), "one override (or None)␣
↳per spec"

for sp in SWEEP_SPECS:
    if not Path(sp).exists():
        warnings.warn(f"Spec not found: {sp}")

print(f"Loading {len(SWEEP_SPECS)} sweep spec(s):")
for sp, od in zip(SWEEP_SPECS, OUTPUT_DIR_OVERRIDES):
    print(f" - {sp} (override={od})")

```

Loading 1 sweep spec(s):

```

- C:\Users\danie\Documents\Python_Scripts\Master Thesis\Thesis-Energy-
Auction\configs\sweeps\default_seeds.yaml
(override=C:\Users\danie\Documents\Python_Scripts\Master Thesis\Thesis-Energy-
Auction\results\sweeps\default)

```

1.4 4. Multi-Sweep Discovery & Loading

```

[5]: from src.utils.run_data import log_exists, default_ep_columns,␣
↳default_yr_columns
from tqdm.auto import tqdm

def _resolve_run_files(variant_name, seed, variant_dir, output_dir):
    """Return (training_log, year_log) paths or (None, None). Tagged then␣
    ↳untagged."""
    search_dirs = []

    # Existing locations
    for d in (variant_dir, output_dir):
        if d is not None:
            d = Path(d)
            if d not in search_dirs:
                search_dirs.append(d)

    # Extra fallback: results/sweep/<tag>[/<variant>]
    # (tag inferred from output_dir name, e.g. default_seeds)
    if output_dir is not None:
        tag = Path(output_dir).name
        alt_base = PROJECT_ROOT / "results" / "sweep" / tag
        for d in (alt_base / variant_name, alt_base):
            if d not in search_dirs:
                search_dirs.append(d)

```

```

training_names = [f"training_log_{variant_name}_s{seed}.csv",
↪f"training_log_s{seed}.csv"]
year_names = [f"year_log_{variant_name}_s{seed}.csv", f"year_log_s{seed}.
↪csv"]

tlog = ylog = None
for base in search_dirs:
    if tlog is None:
        for n in training_names:
            cand = base / n
            if log_exists(cand):
                tlog = cand
                break
    if ylog is None:
        for n in year_names:
            cand = base / n
            if log_exists(cand):
                ylog = cand
                break
    if tlog is not None and ylog is not None:
        break

return tlog, ylog

# Project both logs onto the analysis-superset columns at load time. The
# year_log carries 270+ diagnostic columns; pulling only the ~80 actually
# used by this notebook cuts in-RAM size by ~5-10x. Unknown columns are
# silently dropped by load_run_csv.
EP_COLUMNS = default_ep_columns()
YR_COLUMNS = default_yr_columns()

def _load_csv_or_none(path, columns=None):
    if path is None: return None
    return load_run_csv(path, columns=columns)

def _empty_ep(yr=None):
    if yr is not None and 'episode' in yr.columns:
        return pd.DataFrame({'episode': np.sort(yr['episode'].dropna().
↪unique())})
    return pd.DataFrame({'episode': pd.Series(dtype='int64')})

def _empty_yr(ep=None):
    if ep is not None and 'episode' in ep.columns and len(ep) > 0:
        return pd.DataFrame({'episode': ep['episode'].astype('int64').values,
↪'year': np.zeros(len(ep), dtype='int64')})

```

```

    return pd.DataFrame({'episode': pd.Series(dtype='int64'), 'year': pd.
↳Series(dtype='int64')})

specs = []
for spec_path, override in zip(SWEEP_SPECS, OUTPUT_DIR_OVERRIDES):
    if not Path(spec_path).exists():
        continue
    spec = load_sweep_spec(str(spec_path))
    sweep_name = Path(spec_path).stem
    with open(spec['base_config'], encoding='utf-8') as f:
        base_cfg = yaml.safe_load(f)
    output_dir = Path(override) if override is not None else (PROJECT_ROOT /
↳spec['output_dir']).resolve()
    specs.append({'name': sweep_name, 'spec': spec, 'base_cfg': base_cfg,
↳'output_dir': output_dir})

# Pre-enumerate every (sweep, variant, seed) we'll attempt so tqdm has a real
↳total.
load_jobs = []
for s in specs:
    for v in s['spec']['variants']:
        seeds = v['seeds'] if v['seeds'] is not None else s['spec']['seeds']
        variant_dir = s['output_dir'] / v['name']
        merged_cfg = deep_merge(s['base_cfg'], expand_dotted_overrides(v.
↳get('overrides') or {}))
        flat_overrides = expand_dotted_overrides(v.get('overrides') or {})
        for seed in seeds:
            load_jobs.append((s, v, seed, variant_dir, merged_cfg,
↳flat_overrides))

# Build (sweep, variant, seed) → run dict.
runs = {}
discovery = []
for s, v, seed, variant_dir, merged_cfg, flat_overrides in tqdm(
    load_jobs, desc='Loading sweep CSVs', unit='run', leave=True
):
    tlog, ylog = _resolve_run_files(v['name'], seed, variant_dir,
↳s['output_dir'])
    row = {
        'sweep': s['name'], 'variant': v['name'], 'seed': seed,
        'has_overrides': bool(flat_overrides),
        'training_log': str(tlog) if tlog else None,
        'year_log': str(ylog) if ylog else None,
        'has_training_log': tlog is not None,
        'has_year_log': ylog is not None,
        'found': (tlog is not None) or (ylog is not None),

```

```

    }
    discovery.append(row)
    if not row['found']:
        continue
    ep_df = _load_csv_or_none(tlog, columns=EP_COLUMNS)
    yr_df = _load_csv_or_none(ylog, columns=YR_COLUMNS)
    if ep_df is None and yr_df is None:
        continue
    if ep_df is None: ep_df = _empty_ep(yr_df)
    if yr_df is None: yr_df = _empty_yr(ep_df)
    runs[(s['name'], v['name'], seed)] = {
        'ep_df': ep_df, 'yr_df': yr_df, 'config': merged_cfg,
        'has_overrides': bool(flat_overrides),
        'overrides': flat_overrides,
        'has_training_log': bool(tlog), 'has_year_log': bool(ylog),
    }

discovery_df = pd.DataFrame(discovery)
n_found = int(discovery_df['found'].sum()) if not discovery_df.empty else 0
print(f"Discovered {len(discovery_df)} (sweep, variant, seed) triples; loaded_{n_found}")
if not runs:
    raise RuntimeError("No runs loaded. Check SWEEP_SPECS /_{n_found}")
discovery_df

# Free chunked CSV/parquet readers held by pandas/pyarrow before continuing.
import gc; gc.collect()

```

Loading sweep CSVs: 0% | 0/4 [00:00<?, ?run/s]

Discovered 4 (sweep, variant, seed) triples; loaded 4.

[5]: 27

1.5 5. Default-Only Filter & Sanity Overview

[6]: STRICT_DEFAULT_ONLY = True # if True, drop runs whose variant has any override.

```

if STRICT_DEFAULT_ONLY:
    n_before = len(runs)
    runs = {k: v for k, v in runs.items() if not v['has_overrides']}
    n_after = len(runs)
    if n_after < n_before:
        print(f"Filtered out {n_before - n_after} run(s) whose variant_{n_after} overrides the base config.")
    if n_after == 0:

```



```

        raise RuntimeError("STRICT_DEFAULT_ONLY filtered out every run. Set
        ↳STRICT_DEFAULT_ONLY = False to inspect non-default variants.")

# Treat every surviving run as a default-config seed. Build a flat list.
ALL_RUN_KEYS = sorted(runs.keys())
print(f"\nUsable default-config runs: {len(ALL_RUN_KEYS)}")
for (sw, var, seed), run in sorted(runs.items()):
    n_eps = len(run['ep_df'])
    n_yrs = len(run['yr_df'])
    print(f"    {sw}/{var}/seed={seed:>4}    ({n_eps:>6} episodes, {n_yrs:>7}
    ↳year-rows)")

# Pick a reference config. All defaults agree on schema; just take the first.
REFERENCE_CONFIG = runs[ALL_RUN_KEYS[0]]['config']
N_AGENTS = int(REFERENCE_CONFIG.get('companies', {}).get('n_agents', 8))
N_BOTS = int(REFERENCE_CONFIG.get('companies', {}).get('n_bot_agents', 0))
N_TOTAL = N_AGENTS + N_BOTS
N_YEARS = int(REFERENCE_CONFIG.get('environment', {}).get('n_years', 12))
print(f"\nReference config: n_agents={N_AGENTS}, n_bot_agents={N_BOTS},
↳n_years={N_YEARS}")

# Archetype/reward labelling (default config: 4 archetypes × {financial, ESG}).
ARCHETYPES = ['Coal-heavy', 'Gas-dominant', 'Transitioner', 'Green-leader']
def _agent_archetype(idx0):
    return ARCHETYPES[idx0 // 2] if idx0 // 2 < len(ARCHETYPES) else
    ↳f'Custom-{idx0//2}'
def _agent_weighting(idx0):
    # Even indices (A1, A3, ...) are pure financial; odd are balanced ESG. (Per
    ↳README.)
    return 'Financial' if idx0 % 2 == 0 else 'Balanced ESG'

AGENT_LABELS = [f'A{i+1} ({_agent_archetype(i)}, {_agent_weighting(i)})' for i
↳in range(N_AGENTS)]
print("Agent roster (1-based):")
for lbl in AGENT_LABELS:
    print(f"    {lbl}")

```

Usable default-config runs: 4

```

    default_seeds/default/seed=1729    (120000 episodes, 1440000 year-rows)
    default_seeds/default/seed=5041    (120000 episodes, 1440000 year-rows)
    default_seeds/default/seed=6561    (120000 episodes, 1440000 year-rows)
    default_seeds/default/seed=8191    (120000 episodes, 1440000 year-rows)

```

Reference config: n_agents=8, n_bot_agents=0, n_years=12

Agent roster (1-based):

```

    A1 (Coal-heavy, Financial)

```

A2 (Coal-heavy, Balanced ESG)
A3 (Gas-dominant, Financial)
A4 (Gas-dominant, Balanced ESG)
A5 (Transitioner, Financial)
A6 (Transitioner, Balanced ESG)
A7 (Green-leader, Financial)
A8 (Green-leader, Balanced ESG)

1.5.1 5a. Helpers

A handful of helpers reused across §RQ1–§RQ3. We intentionally treat every loaded run as one seed of the same default-config population (no variant axis) — the helpers therefore aggregate over ALL_RUN_KEYS without any variant grouping.

```
[7]: ROLL = 50                # episode-axis smoothing window
TAIL_EPISODES = 200         # converged-window definition for steady-state metrics

def agent_cols(df, prefix, max_idx=N_AGENTS):
    """Return columns 'prefix_A1' ... 'prefix_A{max_idx}' that exist in df,
    sorted by agent index."""
    rx = re.compile(rf'~{re.escape(prefix)}_A(\d+)$')
    out = []
    for c in df.columns:
        m = rx.match(c)
        if m and int(m.group(1)) <= max_idx:
            out.append((int(m.group(1)), c))
    return [c for _, c in sorted(out)]

def per_run_series(metric_fn, source='ep'):
    """Return dict {run_key: pd.Series indexed by episode} from each run."""
    out = {}
    for k, run in runs.items():
        df = run['yr_df'] if source == 'yr' else run['ep_df']
        if source == 'run':
            s = metric_fn(run)
        else:
            s = metric_fn(df)
        if s is None or len(s) == 0:
            continue
        s = s.copy()
        if source == 'ep' and 'episode' in run['ep_df'].columns:
            s.index = run['ep_df']['episode']
        out[k] = s
    return out

def stack_runs(metric_fn, source='ep'):
    """Return wide DataFrame of episode-indexed series, one column per run."""
    series = per_run_series(metric_fn, source=source)
```

```

    if not series:
        return pd.DataFrame()
    return pd.concat(series, axis=1).sort_index()

def fan_chart(df, ax, color='steelblue', label='runs', smooth=ROLL):
    """Plot median + ±1 std fan + thin per-run lines for a wide DataFrame."""
    if df.empty:
        ax.text(0.5, 0.5, 'no data', ha='center', va='center', transform=ax.
↳transAxes)
        return
    sm = df.rolling(smooth, min_periods=1).mean()
    median = sm.median(axis=1)
    std = sm.std(axis=1)
    for col in sm.columns:
        ax.plot(sm.index, sm[col].values, color=color, lw=0.6, alpha=0.30)
    ax.fill_between(sm.index, (median - std).values, (median + std).values,
        color=color, alpha=0.20, label=f'{label} ±1 std (n={sm.
↳shape[1]})')
    ax.plot(sm.index, median.values, color=color, lw=2.2, label=f'{label}
↳median')
    ax.grid(True, alpha=0.3); ax.legend(fontsize=8, loc='best')

def tail_episodes(ep_df, n=TAIL_EPISODES):
    if ep_df.empty: return ep_df
    return ep_df.tail(min(n, len(ep_df)))

def tail_year_rows(run, n=TAIL_EPISODES):
    ep = tail_episodes(run['ep_df'], n)
    if ep.empty or 'episode' not in run['yr_df'].columns:
        return run['yr_df'].head(0)
    return run['yr_df'][run['yr_df']['episode'].isin(ep['episode'].unique())]

```

1.6 §RQ1 — Stable & Behaviourally Credible Simulation

Three sub-checks:

1. **Convergence** — system reward, composite quality score, and PPO actor/critic losses settle by training end.
2. **Cross-seed reproducibility** — fan charts on headline metrics across loaded seeds; the spread of converged-window means quantifies how reliably the same equilibrium is reached.
3. **Behavioural credibility** — final-window clearing price sits in the EU-ETS-plausible range (~€60–110 per the README’s design notes), compliance rate plateaus high, and warning-flag counts in the converged window are not pathological.

1.6.1 RQ1.1 — Convergence dynamics (rewards, quality score, losses)

```
[8]: def _system_reward(df):
    cols = agent_cols(df, 'reward')
    return df[cols].sum(axis=1) if cols else pd.Series(dtype=float)

def _system_base_reward(df):
    cols = agent_cols(df, 'reward_base')
    return df[cols].sum(axis=1) if cols else pd.Series(dtype=float)

def _quality(df):
    return df['quality_score'] if 'quality_score' in df.columns else pd.
    ↪Series(dtype=float)

def _mean_actor_loss(df):
    cols = agent_cols(df, 'actor_loss')
    return df[cols].mean(axis=1) if cols else pd.Series(dtype=float)

def _mean_critic_loss(df):
    cols = agent_cols(df, 'critic_loss')
    return df[cols].mean(axis=1) if cols else pd.Series(dtype=float)

fig, axes = plt.subplots(2, 2, figsize=(16, 9))
fan_chart(stack_runs(_system_reward), axes[0, 0], color='steelblue',
    ↪label='system reward')
axes[0, 0].set(title='System total reward ( $\Sigma$  reward_A*)', xlabel='Episode',
    ↪ylabel='reward')

fan_chart(stack_runs(_quality), axes[0, 1], color='seagreen',
    ↪label='quality_score')
axes[0, 1].set(title='Composite quality score', xlabel='Episode', ylabel='Q
    ↪[-5, 5]')
axes[0, 1].set_ylim(0, 3)

fan_chart(stack_runs(_mean_actor_loss), axes[1, 0], color='firebrick',
    ↪label='actor loss')
axes[1, 0].set(title='Mean actor loss across agents', xlabel='Episode',
    ↪ylabel='loss')

fan_chart(stack_runs(_mean_critic_loss), axes[1, 1], color='darkorange',
    ↪label='critic loss')
axes[1, 1].set(title='Mean critic loss across agents', xlabel='Episode',
    ↪ylabel='loss')

plt.suptitle('RQ1.1 - Training-time convergence (median  $\pm$ 1 std across loaded
    ↪seeds)',
    fontsize=13, y=1.02)
```

```
plt.tight_layout(); plt.show()
```


C:\Users\danie\AppData\Local\Temp\ipykernel_35472\3832536781.py:36: UserWarning: Glyph 8712 (\N{ELEMENT OF}) missing from font(s) Arial.

```
plt.tight_layout(); plt.show()
```

C:\Users\danie\AppData\Roaming\Python\Python312\site-packages\IPython\core\pylabtools.py:170: UserWarning: Glyph 8712 (\N{ELEMENT OF}) missing from font(s) Arial.

```
fig.canvas.print_figure(bytes_io, **kw)
```

RQ1.1 — Training-time convergence (median ± 1 std across loaded seeds)



1.6.2 RQ1.2 — Cross-seed reproducibility on headline metrics

```
[9]: # Per-seed converged-window means for a small set of headline metrics.
def _converged_metrics_for(run):
    ep = tail_episodes(run['ep_df'])
    yr = tail_year_rows(run)
    out = {}

    rew_cols = agent_cols(ep, 'reward')
    out['system_reward'] = float(ep[rew_cols].sum(axis=1).mean()) if
↪rew_cols else np.nan
    out['quality_score'] = float(ep['quality_score'].mean()) if
↪'quality_score' in ep.columns else np.nan

    if 'clearing_price' in yr.columns:
        out['mean_clearing_price'] = float(yr['clearing_price'].mean())
```

```

elif 'clearing_price_last' in ep.columns:
    out['mean_clearing_price'] = float(ep['clearing_price_last'].mean())
else:
    out['mean_clearing_price'] = np.nan

sf_cols = agent_cols(yr, 'shortfall')
out['compliance_rate'] = float((yr[sf_cols] <= 1e-6).mean(axis=1).mean())
if sf_cols else np.nan

pen_cols = agent_cols(ep, 'penalty')
out['system_penalty'] = float(ep[pen_cols].sum(axis=1).mean()) if pen_cols
else np.nan

inv_cols = agent_cols(ep, 'invest_cost')
out['system_invest'] = float(ep[inv_cols].sum(axis=1).mean()) if inv_cols
else np.nan

gf_cols = agent_cols(ep, 'green_frac')
out['mean_final_green_frac'] = float(ep[gf_cols].mean(axis=1).mean()) if gf_cols
else np.nan

if 'secondary_volume' in ep.columns:
    out['secondary_volume'] = float(ep['secondary_volume'].mean())
else:
    out['secondary_volume'] = np.nan

return out

per_seed_rows = []
for k, run in runs.items():
    row = {'sweep': k[0], 'variant': k[1], 'seed': k[2]}
    row.update(_converged_metrics_for(run))
    per_seed_rows.append(row)
per_seed_df = pd.DataFrame(per_seed_rows)
print(f"Per-seed converged-window means (last {TAIL_EPISODES} episodes):")
display(per_seed_df.round(3))

agg_cols = ['system_reward', 'quality_score', 'mean_clearing_price',
            'compliance_rate',
            'system_penalty', 'system_invest', 'mean_final_green_frac',
            'secondary_volume']
summary = per_seed_df[agg_cols].agg(['mean', 'std', 'min', 'max']).round(3)
summary.loc['cv (%)'] = (summary.loc['std'] / summary.loc['mean'].abs() * 100).
    round(1)
print(f"\nCross-seed summary (n={len(per_seed_df)}):")
summary

```

Per-seed converged-window means (last 200 episodes):

	sweep	variant	seed	system_reward	quality_score	mean_clearing_price	compliance_rate	system_penalty	system_invest	mean_final_green_frac	secondary_volume
0	default_seeds	default	1729	20.853	1.665	66.					
	↪002	0.869	1965.127	5105.939	0.734						
	↪	44.094									
1	default_seeds	default	8191	29.125	1.325	48.					
	↪761	0.850	2750.064	6427.856	0.764						
	↪	40.524									
2	default_seeds	default	6561	12.440	1.326	55.					
	↪902	0.793	3929.627	6847.964	0.770						
	↪	27.174									
3	default_seeds	default	5041	25.155	1.002	46.					
	↪675	0.786	4297.632	5132.392	0.747						
	↪	30.065									

Cross-seed summary (n=4):

	system_reward	quality_score	mean_clearing_price	compliance_rate
system_penalty	system_invest	mean_final_green_frac	secondary_volume	
mean	21.893	1.329	54.335	0.825
3235.613	5878.538	0.754	35.464	
std	7.150	0.270	8.724	0.041
1073.846	893.529	0.017	8.123	
min	12.440	1.002	46.675	0.786
1965.127	5105.939	0.734	27.174	
max	29.125	1.665	66.002	0.869
4297.632	6847.964	0.770	44.094	
cv (%)	32.700	20.300	16.100	5.000
33.200	15.200	2.300	22.900	

Reading the table. A coefficient of variation (cv (%)) below ~10% on the headline outcomes (mean clearing price, compliance rate, mean final green fraction) is the bar we use to call the simulation *reproducible* across seeds — anything significantly larger means the RQ2 / RQ3 claims need to be hedged seed-by-seed.

1.6.3 RQ1.3 — Behavioural credibility in the converged window

```
[10]: # (a) Final-window clearing-price distribution; (b) per-seed compliance plateau.
fig, axes = plt.subplots(1, 2, figsize=(15, 5))

ax = axes[0]
boxes = []
labels = []
for k, run in runs.items():
```

```

yr = tail_year_rows(run)
if 'clearing_price' not in yr.columns or yr.empty:
    continue
boxes.append(yr['clearing_price'].dropna().values)
labels.append(f"{k[0][:8]}/s{k[2]}")
if boxes:
    ax.boxplot(boxes, tick_labels=labels, showfliers=False)
    ax.axhspan(60, 110, color='green', alpha=0.10, label='EU-ETS plausible band_
↳( €60-110/t)')
    ax.set(title='Converged-window clearing price by seed', ylabel='€/tCO ',
↳xlabel='Run')
    ax.legend(fontsize=8); ax.tick_params(axis='x', rotation=30)
else:
    ax.text(0.5, 0.5, 'No clearing_price column in any year_log', ha='center',
↳va='center', transform=ax.transAxes)

ax = axes[1]
for k, run in runs.items():
    sf_cols = agent_cols(run['yr_df'], 'shortfall')
    if not sf_cols or 'episode' not in run['yr_df'].columns:
        continue
    rate = (run['yr_df'][sf_cols] <= 1e-6).mean(axis=1).
↳groupby(run['yr_df']['episode']).mean()
    rate_smooth = rate.rolling(ROLL, min_periods=1).mean()
    ax.plot(rate_smooth.index, rate_smooth.values, lw=1.0, alpha=0.85,
        label=f"{k[0][:8]}/s{k[2]}")
ax.set(title='Compliance rate over training (mean across (agent, year) cells)',
    xlabel='Episode', ylabel='compliance rate [0,1]')
ax.set_ylim(0.0, 1.02); ax.grid(True, alpha=0.3); ax.legend(fontsize=8,
↳loc='best')

plt.suptitle('RQ1.3 - Behavioural credibility: realistic price band +
↳compliance plateau',
    fontsize=13, y=1.02)
plt.tight_layout(); plt.show()

```

C:\Users\danie\AppData\Local\Temp\ipykernel_35472\2922376889.py:36: UserWarning: Glyph 8322 (\N{SUBSCRIPT TWO}) missing from font(s) Arial.

plt.tight_layout(); plt.show()

C:\Users\danie\AppData\Local\Temp\ipykernel_35472\2922376889.py:36: UserWarning: Glyph 8712 (\N{ELEMENT OF}) missing from font(s) Arial.

plt.tight_layout(); plt.show()

C:\Users\danie\AppData\Roaming\Python\Python312\site-packages\IPython\core\pylabtools.py:170: UserWarning: Glyph 8322 (\N{SUBSCRIPT TWO}) missing from font(s) Arial.


fig.canvas.print_figure(bytes_io, **kw)

C:\Users\danie\AppData\Roaming\Python\Python312\site-

packages\IPython\core\pylabtools.py:170: UserWarning: Glyph 8712 (\N{ELEMENT OF}) missing from font(s) Arial.

```
fig.canvas.print_figure(bytes_io, **kw)
```

RQ1.3 — Behavioural credibility: realistic price band + compliance plateau



1.6.4 RQ1.4 — Pathology check: warnings in the converged window

```
[11]: # Sum the warn_* columns across the converged window for each run.
WARN_PREFIX = 'warn_'
warn_rows = []
warn_columns_seen = set()
for k, run in runs.items():
    ep = tail_episodes(run['ep_df'])
    warn_cols = [c for c in ep.columns if c.startswith(WARN_PREFIX)]
    warn_columns_seen.update(warn_cols)
    if not warn_cols:
        continue
    row = {'run': f"{k[0][:8]}/s{k[2]}"}
    # Per-episode mean warning count (interpretable as "how often per episode
    # does this warn fire").
    row.update({c: float(ep[c].mean()) for c in warn_cols})
    warn_rows.append(row)


if warn_rows:
    warn_df = pd.DataFrame(warn_rows).set_index('run').fillna(0.0)
    # Drop columns that are exactly zero across all runs.
    warn_df = warn_df.loc[:, (warn_df > 0).any(axis=0)]
    fig, ax = plt.subplots(figsize=(min(14, 1.5 + 0.45 * len(warn_df.columns)),
    4 + 0.3 * len(warn_df)))
    if warn_df.empty:
        ax.text(0.5, 0.5, 'No non-zero warnings in the converged window - clean
run.',
                ha='center', va='center', transform=ax.transAxes, fontsize=12)
    ax.axis('off')
```

```

else:
    sns.heatmap(warn_df.round(2), annot=True, cmap='Reds', fmt='.2f',
    ↳cbar_kws={'label': 'mean per episode'}, ax=ax)
    ax.set(title=f'Mean warning-flag counts per episode in the converged_
    ↳window (last {TAIL_EPISODES} eps)',
           xlabel='warning detector', ylabel='run')
    plt.xticks(rotation=45, ha='right')
    plt.tight_layout(); plt.show()
    print(f'Inspected {len(warn_columns_seen)} warning columns; showing only_
    ↳the {warn_df.shape[1]} with > 0 events.')
else:
    print('No warn_* columns present in episode logs.')

```

Mean warning-flag counts per episode in the converged window (last 200 eps)



Inspected 6 warning columns; showing only the 4 with > 0 events.

RQ1 summary. The convergence panels (1.1) plus the cross-seed table (1.2) and credibility plots (1.3 / 1.4) together provide the evidence base for Sub-RQ 1: with the default config, the simulation converges, reproduces a plausible EU-ETS price band and compliance plateau across loaded seeds,

and the converged window is not dominated by pathology warnings.

1.7 §RQ2 — Regulatory-Mechanism Behaviour Within the Default Config

Scope caveat. Real *sensitivity* analysis (cap LRF / MSR on/off / penalty level) requires sweep variants and is deferred to a separate experiments notebook. What we *can* show with default-config data is a descriptive characterisation of how the regulatory mechanisms operate in equilibrium: how the cap maps to the clearing price, how the MSR withholds and releases, when the penalty actually binds, and how quickly the system decarbonises. This is the baseline that the sensitivity analysis will perturb.

Four panels:

- **2.1** — Cap → price linkage by year (mean across episodes/seeds in the converged window).
- **2.2** — TNAC trajectory and MSR reserve / flow behaviour.
- **2.3** — Penalty incidence: how often does it bind, and how big is it when it does?
- **2.4** — Decarbonisation pace: green-share trajectory and MAC-driven reductions.

1.7.1 RQ2.1 — Cap and clearing-price trajectory by year

```
[12]: # Average across all seeds and converged-window episodes.
year_frames = []
for k, run in runs.items():
    yr = tail_year_rows(run)
    if yr.empty: continue
    keep = ['episode', 'year', 'cap', 'clearing_price', 'auction_volume',
    ↪ 'tnac',
            'msr_reserve', 'msr_withhold_this_year', 'msr_release_this_year',
            'msr_total_cancelled', 'secondary_price']
    keep = [c for c in keep if c in yr.columns]
    if 'year' not in keep: continue
    df = yr[keep].copy(); df['__seed__'] = k[2]
    year_frames.append(df)
if not year_frames:
    raise RuntimeError('No year_log data available - RQ2 requires year-level_
    ↪ logs.')
yr_pool = pd.concat(year_frames, ignore_index=True)

by_year = yr_pool.groupby('year').agg(['mean', 'std'])

fig, axes = plt.subplots(1, 2, figsize=(14, 4.5))

ax = axes[0]
if ('cap', 'mean') in by_year.columns:
    ax.plot(by_year.index, by_year[('cap', 'mean')], lw=2, color='black',
    ↪ label='cap (Mt, mean)')
    ax.set_ylabel('cap (Mt)', color='black')
```

```


ax2 = ax.twinx()
if ('clearing_price', 'mean') in by_year.columns:
    ax2.plot(by_year.index, by_year[('clearing_price', 'mean')], lw=2,
    color='firebrick',
    label='clearing price (€/t)')
    ax2.fill_between(by_year.index,
                    by_year[('clearing_price', 'mean')] -
    by_year[('clearing_price', 'std')],
                    by_year[('clearing_price', 'mean')] +
    by_year[('clearing_price', 'std')],
                    color='firebrick', alpha=0.15)
    ax2.set_ylabel('clearing price (€/t)', color='firebrick')
ax.set(title='Cap vs clearing price by simulation year (converged window)',
xlabel='year')
ax.grid(True, alpha=0.3)

ax = axes[1]
if ('auction_volume', 'mean') in by_year.columns:
    ax.plot(by_year.index, by_year[('auction_volume', 'mean')], lw=2,
    color='steelblue', label='auction volume (Mt)')
if ('cap', 'mean') in by_year.columns:
    ax.plot(by_year.index, by_year[('cap', 'mean')], lw=1.5, color='black',
    linestyle='--', label='cap (Mt)')
ax.set(title='Cap vs realised auction volume by year (MSR + rollover impact)',
    xlabel='year', ylabel='Mt')
ax.legend(fontsize=9); ax.grid(True, alpha=0.3)

plt.suptitle('RQ2.1 - Regulatory cap schedule and its realised market impact',
    fontsize=13, y=1.02)
plt.tight_layout(); plt.show()

```

RQ2.1 — Regulatory cap schedule and its realised market impact



1.7.2 RQ2.2 — TNAC, MSR reserve, and MSR flow


```
[13]: fig, axes = plt.subplots(1, 3, figsize=(18, 4.5))

ax = axes[0]
if ('tnac', 'mean') in by_year.columns:
    ax.plot(by_year.index, by_year[('tnac', 'mean')], lw=2, color='darkblue')
    ax.fill_between(by_year.index,
                    by_year[('tnac', 'mean')] - by_year[('tnac', 'std')],
                    by_year[('tnac', 'mean')] + by_year[('tnac', 'std')],
                    color='darkblue', alpha=0.15)
    ax.set(title='TNAC trajectory (Total Number of Allowances in Circulation)',
           xlabel='year', ylabel='Mt')
    ax.grid(True, alpha=0.3)

ax = axes[1]
if ('msr_reserve', 'mean') in by_year.columns:
    ax.plot(by_year.index, by_year[('msr_reserve', 'mean')], lw=2,
            color='purple', label='MSR reserve')
if ('msr_total_cancelled', 'mean') in by_year.columns:
    ax.plot(by_year.index, by_year[('msr_total_cancelled', 'mean')], lw=2,
            color='red',
            linestyle='--', label='cumulative cancelled')
ax.set(title='MSR reserve & cancellations by year', xlabel='year', ylabel='Mt')
ax.legend(fontsize=9); ax.grid(True, alpha=0.3)

ax = axes[2]
years = by_year.index.values
if ('msr_withhold_this_year', 'mean') in by_year.columns:
    wh = by_year[('msr_withhold_this_year', 'mean')]
    ax.bar(years - 0.18, wh, width=0.35, color='darkorange', label='withhold')
if ('msr_release_this_year', 'mean') in by_year.columns:
    rel = by_year[('msr_release_this_year', 'mean')]
    ax.bar(years + 0.18, rel, width=0.35, color='seagreen', label='release')
ax.axhline(0, color='black', lw=0.8)
ax.set(title='MSR yearly flow', xlabel='year', ylabel='Mt')
ax.legend(fontsize=9); ax.grid(True, alpha=0.3)

plt.suptitle('RQ2.2 - Market Stability Reserve dynamics in the default config',
             fontsize=13, y=1.02)
plt.tight_layout(); plt.show()
```



1.7.3 RQ2.3 — Penalty incidence and magnitude

```
[14]: # (a) Share of (agent, year) cells with shortfall > 0; (b) mean per-cell
      ↪penalty when it bites.
share_rows = []
for k, run in runs.items():
    yr = tail_year_rows(run)
    sf_cols = agent_cols(yr, 'shortfall')
    pen_cols = agent_cols(yr, 'penalty')
    if not sf_cols: continue
    sf = yr[sf_cols].values.flatten()
    pen = yr[pen_cols].values.flatten() if pen_cols else None
    nc_mask = sf > 1e-6
    share_rows.append({
        'run': f'{k[0][:8]}/s{k[2]}',
        'pct_cells_short': 100.0 * nc_mask.mean() if sf.size else np.nan,
        'mean_shortfall_when_short_Mt': float(sf[nc_mask].mean()) if nc_mask.
        ↪any() else 0.0,
        'mean_penalty_when_short_Meur': float(pen[nc_mask].mean()) if (pen is
        ↪not None and nc_mask.any()) else 0.0,
    })
share_df = pd.DataFrame(share_rows)
print(f'Penalty incidence in the converged window (last {TAIL_EPISODES}
      ↪episodes):')
display(share_df.round(3))

# UDBC mix at the system level (averaged across seeds).
ud_rows = []
for k, run in runs.items():
    ep = tail_episodes(run['ep_df'])
    for b in ['U', 'D', 'M', 'B', 'C']:
        cols = agent_cols(ep, f'udbc_{b}_total')
        if cols:
            ud_rows.append({'run': f'{k[0][:8]}/s{k[2]}', 'bucket': b,
```


```

        'mean_total_per_episode': float(ep[cols].
↪sum(axis=1).mean()))}
if ud_rows:
    ud_df = pd.DataFrame(ud_rows).pivot(index='run', columns='bucket',
↪values='mean_total_per_episode').fillna(0.0)
    ud_df = ud_df[[c for c in ['U', 'D', 'M', 'B', 'C'] if c in ud_df.columns]]
    fig, ax = plt.subplots(figsize=(10, 4))
    ud_df.plot(kind='bar', stacked=True, ax=ax,
        color=['#d62728', '#9467bd', '#ff7f0e', '#2ca02c', '#1f77b4'])
    ax.set(title='UDBC compliance-pathway counts per episode (converged
↪window)',
        ylabel='# (agent·year) cells per episode', xlabel='run')
    ax.legend(title='bucket', fontsize=8); ax.tick_params(axis='x', rotation=30)
    plt.tight_layout(); plt.show()
    print('U = pure under-bid; D = pure debt-cascade; M = mixed; B =
↪bank-covered; C = secondary-covered.')

```

Penalty incidence in the converged window (last 200 episodes):

	run	pct_cells_short	mean_shortfall_when_short_Mt	↪mean_penalty_when_short_Meur
0	default_/s1729	13.094	0.976	↪156.335
1	default_/s8191	15.010	1.188	↪190.844
2	default_/s6561	20.677	1.251	↪197.966
3	default_/s5041	21.411	1.329	↪209.080



U = pure under-bid; D = pure debt-cascade; M = mixed; B = bank-covered; C = secondary-covered.

1.7.4 RQ2.4 — Decarbonisation pace

```
[15]: # (a) Mean green_frac by year across episodes/seeds; (b) MAC reduction by year.
yr_pool2 = []
for k, run in runs.items():
    yr = tail_year_rows(run)
    keep = ['year']
    gf_cols = agent_cols(yr, 'green_frac')
    mac_cols = agent_cols(yr, 'mac_reduction')
    em_cols = agent_cols(yr, 'emissions')
    if 'year' not in yr.columns: continue
    df = yr[['year']].copy()
    if gf_cols: df['green_frac_mean'] = yr[gf_cols].mean(axis=1)
    if mac_cols: df['mac_reduction_total'] = yr[mac_cols].sum(axis=1)
    if em_cols: df['emissions_total'] = yr[em_cols].sum(axis=1)
    yr_pool2.append(df)
if yr_pool2:
    pool = pd.concat(yr_pool2, ignore_index=True).groupby('year').agg(['mean',
    ↪ 'std'])

    fig, axes = plt.subplots(1, 3, figsize=(18, 4.5))
    ax = axes[0]
    if ('green_frac_mean', 'mean') in pool.columns:
        ax.plot(pool.index, pool[('green_frac_mean', 'mean')], lw=2,
    ↪ color='seagreen')
        ax.fill_between(pool.index,
                        pool[('green_frac_mean', 'mean')] -
    ↪ pool[('green_frac_mean', 'std')],
                        pool[('green_frac_mean', 'mean')] +
    ↪ pool[('green_frac_mean', 'std')],
                        color='seagreen', alpha=0.15)
        ax.set(title='Mean green generation share by year', xlabel='year',
    ↪ ylabel='green frac [0,1]')
        ax.set_ylim(0, 1); ax.grid(True, alpha=0.3)

    ax = axes[1]
    if ('emissions_total', 'mean') in pool.columns:
        ax.plot(pool.index, pool[('emissions_total', 'mean')], lw=2,
    ↪ color='black')
        ax.set(title='System total emissions by year', xlabel='year',
    ↪ ylabel='Mt')
        ax.grid(True, alpha=0.3)

    ax = axes[2]
    if ('mac_reduction_total', 'mean') in pool.columns:
        ax.bar(pool.index, pool[('mac_reduction_total', 'mean')],
    ↪ color='steelblue')
```



```

ax.set(title='MAC fuel-switching reductions by year', xlabel='year',
       ylabel='Mt abated')
ax.grid(True, alpha=0.3)

plt.suptitle('RQ2.4 - Decarbonisation pace within the default config',
             fontsize=13, y=1.02)
plt.tight_layout(); plt.show()


```

C:\Users\danie\AppData\Local\Temp\ipykernel_35472\1425922483.py:43: UserWarning: Glyph 8712 (\N{ELEMENT OF}) missing from font(s) Arial.

```
plt.tight_layout(); plt.show()
```

C:\Users\danie\AppData\Roaming\Python\Python312\site-packages\IPython\core\pylabtools.py:170: UserWarning: Glyph 8712 (\N{ELEMENT OF}) missing from font(s) Arial.

```
fig.canvas.print_figure(bytes_io, **kw)
```



RQ2 summary. Within the default config, RQ2.1–2.4 give the *descriptive* baseline of how cap, MSR, penalty, and emission outcomes move together. The sensitivity question — *would these trajectories look qualitatively different under a tighter LRF, or with the MSR off?* — is what the cross-config experiments notebook will answer.

1.8 §RQ3 — Strategic Behaviour and Aggregation

The default config crosses 4 archetypes (coal-heavy, gas-dominant, transitioner, green-leader) with 2 reward weightings (pure financial, balanced ESG). That gives us the financial-vs-environmental contrast that RQ3 asks about, *inside one config*.

Five panels:

- **3.1** — Bidding behaviour by archetype × weighting.
- **3.2** — Investment behaviour by archetype × weighting.
- **3.3** — Secondary-market role by archetype.
- **3.4** — UDBC compliance pathway by archetype.
- **3.5** — Reward decomposition (financial vs ESG vs penalty) by weighting.

1.8.1 RQ3.1 — Bidding behaviour by archetype × reward-weighting

```
[16]: # Per-agent converged-window means, then group by (archetype, weighting).
def _per_agent_table(run, episode_cols, year_cols):
    """Return list of dicts (one per agent) with mean values from converged_
    ↪window."""
    ep = tail_episodes(run['ep_df'])
    yr = tail_year_rows(run)
    rows = []
    for i in range(N_AGENTS):
        row = {'agent_idx': i, 'archetype': _agent_archetype(i), 'weighting':_
        ↪_agent_weighting(i)}
        for prefix in episode_cols:
            col = f'{prefix}_A{i+1}'
            row[prefix] = float(ep[col].mean()) if col in ep.columns else np.nan
        for prefix in year_cols:
            col = f'{prefix}_A{i+1}'
            row[prefix] = float(yr[col].mean()) if col in yr.columns else np.nan
        rows.append(row)
    return rows

agent_rows = []
for k, run in runs.items():
    for r in _per_agent_table(
        run,
        episode_cols=['bid_price', 'avg_bid_mult', 'avg_bid_coverage',
        ↪'sec_buy_vol', 'sec_sell_vol', 'sec_buy_avg_px',_
        ↪'sec_sell_avg_px',
        ↪'sec_buy_years', 'sec_sell_years',
        ↪'invest_cost', 'green_frac', 'delta_green',
        ↪'mean_alloc', 'penalty', 'shortfall', 'reward',_
        ↪'reward_base',
        ↪'inv_onshore_share', 'inv_offshore_share',_
        ↪'inv_solar_share'],
        year_cols=['mac_reduction', 'mac_cost'],
    ):
        r.update({'sweep': k[0], 'seed': k[2]})
        agent_rows.append(r)
agent_df = pd.DataFrame(agent_rows)
print(f'Per-(run, agent) rows: {len(agent_df)}')

fig, axes = plt.subplots(1, 3, figsize=(18, 4.5))
order = [(arch, w) for arch in ARCHETYPES for w in ['Financial', 'Balanced_
    ↪ESG']]
labels = [f'{a}\n{w}' for a, w in order]

def _bar(ax, col, title, ylabel):
```

```

means = []; stds = []
for arch, w in order:
    sub = agent_df[(agent_df['archetype'] == arch) & (agent_df['weighting']_
    == w)][col]
    means.append(sub.mean()); stds.append(sub.std())
x = np.arange(len(order))
colors = ['#d62728' if w == 'Financial' else '#2ca02c' for _, w in order]
ax.bar(x, means, yerr=stds, color=colors, capsize=3, alpha=0.85)
ax.set_xticks(x); ax.set_xticklabels(labels, fontsize=8, rotation=30,
    ha='right')
ax.set(title=title, ylabel=ylabel); ax.grid(True, alpha=0.3, axis='y')

_bar(axes[0], 'bid_price', 'Mean Phase-1 bid price', '€/tCO2 ')
_bar(axes[1], 'avg_bid_mult', 'Mean qty multiplier on `estimate_need`,
    'multiplier')
_bar(axes[2], 'avg_bid_coverage', 'Mean coverage = alloc / need', 'ratio
    (clipped to 1)')
plt.suptitle('RQ3.1 - Bidding behaviour: red = pure financial reward; green =
    balanced ESG',
    fontsize=12, y=1.04)
plt.tight_layout(); plt.show()

```


Per-(run, agent) rows: 32

C:\Users\danie\AppData\Local\Temp\ipykernel_35472\309587121.py:55: UserWarning: Glyph 8322 (\N{SUBSCRIPT TWO}) missing from font(s) Arial.

```
plt.tight_layout(); plt.show()
```

C:\Users\danie\AppData\Roaming\Python\Python312\site-packages\IPython\core\pylabtools.py:170: UserWarning: Glyph 8322 (\N{SUBSCRIPT TWO}) missing from font(s) Arial.

```
fig.canvas.print_figure(bytes_io, **kw)
```



1.8.2 RQ3.2 — Investment behaviour

```
[17]: fig, axes = plt.subplots(1, 3, figsize=(18, 4.5))
_bar(axes[0], 'invest_cost', 'Mean episode invest cost', 'M€')
_bar(axes[1], 'green_frac', 'End-of-episode green fraction', 'green frac [0,1]')
_bar(axes[2], 'delta_green', 'Δ green fraction over the episode', 'Δ green frac')

plt.suptitle('RQ3.2 - Investment outcomes: red = pure financial; green = balanced ESG',
            fontsize=12, y=1.04)
plt.tight_layout(); plt.show()


# Tech mix by archetype × weighting (stacked bar).
tech_means = (agent_df.groupby(['archetype', 'weighting'])
              [['inv_onshore_share', 'inv_offshore_share', 'inv_solar_share']].
              mean())
fig, ax = plt.subplots(figsize=(10, 4.5))
tech_means.plot(kind='bar', stacked=True, ax=ax,
               color=['#1f77b4', '#17becf', '#ffbf00'])
ax.set(title='Tech mix of episode investments by archetype × weighting',
      ylabel='share of invested MW', xlabel='')
ax.legend(['onshore wind', 'offshore wind', 'solar'], fontsize=9)
plt.xticks(rotation=30, ha='right')
plt.tight_layout(); plt.show()
```


C:\Users\danie\AppData\Local\Temp\ipykernel_35472\1602023298.py:7: UserWarning: Glyph 8712 (\N{ELEMENT OF}) missing from font(s) Arial.

```
plt.tight_layout(); plt.show()
```

C:\Users\danie\AppData\Roaming\Python\Python312\site-packages\IPython\core\pylabtools.py:170: UserWarning: Glyph 8712 (\N{ELEMENT OF}) missing from font(s) Arial.

```
fig.canvas.print_figure(bytes_io, **kw)
```





1.8.3 RQ3.3 — Secondary-market role

```
[18]: fig, axes = plt.subplots(1, 3, figsize=(18, 4.5))
_bar(axes[0], 'sec_buy_vol', 'Mean Mt bought on secondary', 'Mt')
_bar(axes[1], 'sec_sell_vol', 'Mean Mt sold on secondary', 'Mt')
# Net buyer/seller years (out of 12).
def _net_buyer_years(row):
    by = row.get('sec_buy_years', np.nan); sy = row.get('sec_sell_years', np.
    ↪nan)
    if pd.isna(by) or pd.isna(sy): return np.nan
    return by - sy
agent_df['net_buyer_years'] = agent_df.apply(_net_buyer_years, axis=1)
_bar(axes[2], 'net_buyer_years', '(buy - sell) years per episode', 'years')
plt.suptitle('RQ3.3 - Who buys, who sells: secondary-market role by archetype ×_
    ↪weighting',
            fontsize=12, y=1.04)
plt.tight_layout(); plt.show()
```




1.8.4 RQ3.4 — Compliance pathway (UDBC) by archetype × weighting

```
[19]: # Mean per-agent UDBC totals + grouped by archetype × weighting + stacked bar
      ↪ of shares.
udbc_buckets = ['U', 'D', 'M', 'B', 'C']
udbc_rows = []
for k, run in runs.items():
    ep = tail_episodes(run['ep_df'])
    for i in range(N_AGENTS):
        row = {'archetype': _agent_archetype(i), 'weighting':
            ↪ _agent_weighting(i)}
        for b in udbc_buckets:
            col = f'udbc_{b}_total_A{i+1}'
            row[b] = float(ep[col].mean()) if col in ep.columns else 0.0
        udbc_rows.append(row)
udbc_df = pd.DataFrame(udbc_rows)
udbc_grouped = udbc_df.groupby(['archetype', 'weighting'])[udbc_buckets].mean()
udbc_share = udbc_grouped.div(udbc_grouped.sum(axis=1).replace(0, np.nan),
    ↪ axis=0).fillna(0.0)

fig, axes = plt.subplots(1, 2, figsize=(15, 4.5))
udbc_grouped.plot(kind='bar', stacked=True, ax=axes[0],
    color=['#d62728', '#9467bd', '#ff7f0e', '#2ca02c', '#1f77b4'])
axes[0].set(title='UDBC counts per episode by archetype × weighting',
    ylabel='# (agent·year) cells per episode', xlabel='')
axes[0].legend(title='bucket', fontsize=8)
axes[0].tick_params(axis='x', rotation=30)

udbc_share.plot(kind='bar', stacked=True, ax=axes[1],
    color=['#d62728', '#9467bd', '#ff7f0e', '#2ca02c', '#1f77b4'])
axes[1].set(title='UDBC share-within-row by archetype × weighting',
    ylabel='share', xlabel='')
axes[1].set_ylim(0, 1.02)
axes[1].legend(title='bucket', fontsize=8)
axes[1].tick_params(axis='x', rotation=30)
plt.suptitle('RQ3.4 - Compliance pathway by archetype × reward-weighting',
    fontsize=12, y=1.04)
plt.tight_layout(); plt.show()
```



1.8.5 RQ3.5 — Reward decomposition: financial vs balanced-ESG agents

```
[20]: # At the (run, agent) level we have reward_base = financial - penalty + ESG
      ↪ (pre-shaping).
# We don't have the financial / ESG / penalty buckets cleanly split in the
      ↪ episode CSV (they're
# combined into reward_base), so we use what IS exposed: penalty_A{i},
      ↪ invest_cost_A{i},
# reward_A{i}, reward_base_A{i}.

decomp = agent_df.groupby(['archetype', 'weighting'])[['reward', 'reward_base',
      ↪ 'penalty', 'invest_cost']].mean()
decomp['shaping'] = decomp['reward'] - decomp['reward_base']
decomp_print = decomp[['reward', 'reward_base', 'shaping', 'penalty',
      ↪ 'invest_cost']].round(2)
print('Per-archetype × weighting mean reward decomposition (M€-scale,
      ↪ post-normalisation units):')
display(decomp_print)

fig, ax = plt.subplots(figsize=(11, 4.5))
order = [(a, w) for a in ARCHETYPES for w in ['Financial', 'Balanced ESG']]
labels = [f'{a}\n{w}' for a, w in order]
x = np.arange(len(order))
width = 0.22

rew_b   = [decomp.loc[(a, w), 'reward_base'] for a, w in order]
shape_  = [decomp.loc[(a, w), 'shaping']     for a, w in order]
pen     = [decomp.loc[(a, w), 'penalty']      for a, w in order]
inv     = [decomp.loc[(a, w), 'invest_cost']  for a, w in order]

ax.bar(x - 1.5 * width, rew_b, width, label='reward_base',
      ↪ (financial+ESG-penalty)', color='steelblue')
```


```

ax.bar(x - 0.5 * width, shape_, width, label='shaping (cov-gap, opp-cost,
    ↪green, terminal)', color='seagreen')
ax.bar(x + 0.5 * width, -np.asarray(pen), width, label='-penalty (M€)',
    ↪color='firebrick')
ax.bar(x + 1.5 * width, -np.asarray(inv), width, label='-invest_cost (M€)',
    ↪color='darkorange')
ax.axhline(0, color='black', lw=0.8)
ax.set_xticks(x); ax.set_xticklabels(labels, fontsize=8, rotation=30,
    ↪ha='right')
ax.set(title='RQ3.5 - Decomposition of converged-window mean reward by
    ↪archetype × weighting',
        ylabel='component value (per episode)')
ax.legend(fontsize=9, loc='best'); ax.grid(True, alpha=0.3, axis='y')
plt.tight_layout(); plt.show()

```

Per-archetype × weighting mean reward decomposition (M€-scale, post-normalisation units):

		reward	reward_base	shaping	penalty	invest_cost
Coal-heavy	Balanced ESG	-1.94	-2.45	0.51	239.98	945.64
	Financial	-7.71	-8.19	0.47	1561.00	681.41
Gas-dominant	Balanced ESG	5.18	4.66	0.52	180.35	1062.66
	Financial	-10.09	-10.58	0.49	859.45	391.04
Green-leader	Balanced ESG	10.48	9.95	0.53	15.94	893.53
	Financial	10.77	10.24	0.53	50.80	303.47
Transitioner	Balanced ESG	5.11	4.59	0.52	218.20	774.03
	Financial	10.11	9.58	0.52	109.90	826.76



Aggregation note. The system-level outcomes that RQ3 asks about — total clearing price, system green fraction, total penalty volume — are exactly what RQ1.2 / RQ1.3 / RQ2.1 already report. The missing piece RQ3 adds is *attribution*: which archetypes drive those aggregates? The \$RQ3

panels above answer that — and in particular, panels 3.1 / 3.4 / 3.5 contrast the pure-financial agent (red) against its balanced-ESG twin (green) within each archetype, which is the direct test of “how do financial vs environmental objectives shape strategy”.

1.8.6 RQ4 — Compliance pressure, debt cascade, and credit stress

These panels read from the v8.6.1 expanded year-log columns. Cells silently no-op on logs from older versions (every column access is guarded by `if col in df.columns`). The four sub-questions are:

- **RQ4.1** — Does coverage gap ($\max(0, \text{need} - \text{alloc})$ Mt) shrink as agents learn to size their auction bids? Is the residual gap eaten by the secondary market or does it propagate as carry-forward debt?
- **RQ4.2** — Auction concentration: does the system stabilise around a diffuse allocation, or does a small subset of bidders capture most volume (HHI / max-agent-share)?
- **RQ4.3** — Secondary-market liquidity: how many agents participate on each side, and what fraction of intent gets executed?
- **RQ4.4** — Credit stress in the converged window: peak emergency-loan balance and peak carry-forward debt by archetype $\times$ reward-weighting.

1.8.7 RQ4.1 — Coverage gap & carry-forward trajectory

```
[21]: # RQ4.1 - Per-year coverage gap and end-of-year carry-forward (debt cascade).
# Cross-seed median + ±1 std fan, computed from the year_log only - works on
# any v8.6.1+ run. Earlier logs silently produce empty figures.
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

def _stack_year_metric(prefix, max_idx=N_TOTAL, agg="sum"):
    """Return wide DataFrame indexed by year with one column per run, where
    each value is the cross-agent (sum/mean) of `{prefix}_A{i}` for that
    year - averaged across episodes within the run."""
    out = {}
    for k, run in runs.items():
        yr = run["yr_df"]
        if "year" not in yr.columns:
            continue
        cols = [f"{prefix}_A{i+1}" for i in range(max_idx)
                if f"{prefix}_A{i+1}" in yr.columns]
        if not cols:
            continue
        # Per-year mean across episodes of the cross-agent agg.
        if agg == "sum":
            per_year = yr.groupby("year")[cols].mean().sum(axis=1)
        else:
            per_year = yr.groupby("year")[cols].mean().mean(axis=1)
```


```

        out[k] = per_year
    return pd.concat(out, axis=1).sort_index() if out else pd.DataFrame()

cov_gap_df = _stack_year_metric("coverage_gap")           # system-sum (Mt)
cf_end_df = _stack_year_metric("carry_forward_end")       # system-sum (Mt)
cf_start_df = _stack_year_metric("carry_forward_start")   # system-sum (Mt)

if cov_gap_df.empty and cf_end_df.empty:
    print("[RQ4.1] coverage_gap_A* / carry_forward_*_A* columns missing - pre-8.
    ↪6.1 log.")
else:
    fig, axes = plt.subplots(1, 2, figsize=(13, 4))
    for ax, df, ttl in [
        (axes[0], cov_gap_df, "System coverage gap (Mt) - pre-secondary"),
        (axes[1], cf_end_df, "System end-of-year carry-forward (Mt)"),
    ]:
        if df.empty:
            ax.set_title(ttl + "\n(missing)"); continue
        med = df.median(axis=1); std = df.std(axis=1)
        ax.fill_between(med.index, med - std, med + std, alpha=0.2,
        ↪color="steelblue")
        ax.plot(med.index, med, "o-", color="steelblue", label="median across
        ↪seeds")
        ax.set_xlabel("Year"); ax.set_ylabel("Mt"); ax.set_title(ttl)
        ax.grid(alpha=0.3); ax.legend(loc="best", fontsize=9)
    plt.tight_layout(); plt.show()

```



1.8.8 RQ4.2 — Auction concentration (HHI, max-agent share)

```

[22]: # RQ4.2 - Auction concentration trajectory. Reads `auction_hhi` and
      # `auction_max_agent_share` (year-level scalars added in v8.6.1).
      import numpy as np
      import pandas as pd
      import matplotlib.pyplot as plt

```

```


def _stack_year_scalar(col):
    out = {}
    for k, run in runs.items():
        yr = run["yr_df"]
        if "year" not in yr.columns or col not in yr.columns:
            continue
        out[k] = yr.groupby("year")[col].mean()
    return pd.concat(out, axis=1).sort_index() if out else pd.DataFrame()

hhi_df = _stack_year_scalar("auction_hhi")
mas_df = _stack_year_scalar("auction_max_agent_share")
fail_df = _stack_year_scalar("auction_failed")
unsold_df = _stack_year_scalar("auction_unsold")

if hhi_df.empty and mas_df.empty:
    print("RQ4.2 auction_hhi / auction_max_agent_share missing - pre-8.6.1_
↳log.")
else:
    fig, axes = plt.subplots(1, 3, figsize=(15, 4))
    for ax, df, ttl, ylabel in [
        (axes[0], hhi_df, "Auction HHI (allocation concentration)", "HHI"),
        (axes[1], mas_df, "Max single-agent share", "Share"),
        (axes[2], unsold_df, "Auction unsold (Mt)", "Mt"),
    ]:
        if df.empty:
            ax.set_title(ttl + "\n(missing)"); continue
        med = df.median(axis=1); std = df.std(axis=1)
        ax.fill_between(med.index, med - std, med + std, alpha=0.2,
↳color="darkorange")
        ax.plot(med.index, med, "o-", color="darkorange")
        ax.set_xlabel("Year"); ax.set_ylabel(ylabel); ax.set_title(ttl)
        ax.grid(alpha=0.3)
    plt.tight_layout(); plt.show()

    if not fail_df.empty:
        n_failures = int(np.nansum(fail_df.values))
        print(f"Total auction-failure (year, seed) cells across loaded runs:
↳{n_failures}")

```



Total auction-failure (year, seed) cells across loaded runs: 0

1.8.9 RQ4.3 — Secondary-market liquidity (intent vs execution)

```
[23]: # RQ4.3 - Secondary-market participation: how many agents posted intent
# vs how many actually traded? A wide intent-execution gap means agents
# repeatedly try to clear at uncrossable prices.
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

cols = {
    "buyers_intent": "secondary_n_buyers_intent",
    "sellers_intent": "secondary_n_sellers_intent",
    "buyers_executed": "secondary_n_buyers_executed",
    "sellers_executed": "secondary_n_sellers_executed",
}


present = {label: col for label, col in cols.items() if any(col in run["yr_df"].
    ↪columns for run in runs.values())}

if not present:
    print("[RQ4.3] secondary_n_*_intent/executed columns missing - pre-8.6.1_
    ↪log.")
else:
    fig, ax = plt.subplots(figsize=(11, 4))
    for label, col in present.items():
        rows = []
        for k, run in runs.items():
            yr = run["yr_df"]
            if col not in yr.columns or "year" not in yr.columns:
                continue
            rows.append(yr.groupby("year")[col].mean())
        if not rows:
            continue
        s = pd.concat(rows, axis=1).median(axis=1)
        ls = "--" if "executed" in label else "--"
```

```

ax.plot(s.index, s, ls + "o", label=label)
ax.set_xlabel("Year"); ax.set_ylabel("Number of agents")
ax.set_title("Secondary-market participation - intent vs executed")
ax.legend(); ax.grid(alpha=0.3)
plt.tight_layout(); plt.show()

```



1.8.10 RQ4.4 — Converged-window credit stress by archetype × weighting

```

[24]: # RQ4.4 - Per-(archetype, weighting) cell of the mean peak emergency-loan
# balance and peak carry-forward debt in the converged window. Sourced from
# the per-episode `peak_loan_outstanding_A{i}` / `peak_carry_forward_A{i}`
# (v8.6.1+); pre-8.6.1 logs lack these columns and produce an empty table.
import numpy as np
import pandas as pd

def _converged_credit_metrics():
    rows = []
    for k, run in runs.items():
        ep_tail = tail_episodes(run["ep_df"])
        cfg = run["config"]
        for i in range(N_AGENTS):
            arche = ARCHETYPES[i // 2] if i // 2 < len(ARCHETYPES) else
↪ f"A{i+1}"
            w_cost, w_green = (cfg.get("companies", {}).get("reward_weights")
↪ or [[1, 0]] * N_AGENTS)[i]
            weighting = "Financial" if w_green < 0.25 else ("ESG" if w_green >
↪ 0.75 else "Balanced")
            row = {"sweep": k[0], "variant": k[1], "seed": k[2],
                    "archetype": arche, "weighting": weighting, "agent": i + 1}
            for stub in ["peak_loan_outstanding", "peak_carry_forward",
                        "final_treasury_reserve"]:
                col = f"{stub}_A{i+1}"

```

```

        row[stub] = float(ep_tail[col].mean()) if col in ep_tail.
    columns else np.nan
    rows.append(row)
    return pd.DataFrame(rows)

credit_df = _converged_credit_metrics()
if credit_df.empty or credit_df[["peak_loan_outstanding",
    "peak_carry_forward"]].isna().all().all():
    print("[RQ4.4] peak_loan_outstanding_A* / peak_carry_forward_A* missing -
    pre-8.6.1 log.")
else:
    pivot_loan = credit_df.pivot_table(
        index="archetype", columns="weighting", values="peak_loan_outstanding",
        aggfunc="mean"
    ).round(3)
    pivot_cf = credit_df.pivot_table(
        index="archetype", columns="weighting", values="peak_carry_forward",
        aggfunc="mean"
    ).round(3)
    print("Mean peak emergency-loan balance (M€) - converged window:")
    display(pivot_loan)
    print("\nMean peak end-of-year carry-forward (Mt) - converged window:")
    display(pivot_cf)

```

Mean peak emergency-loan balance (M€) - converged window:

weighting archetype	Balanced	Financial
Coal-heavy	0.0	0.0
Gas-dominant	0.0	0.0
Green-leader	0.0	0.0
Transitioner	0.0	0.0

Mean peak end-of-year carry-forward (Mt) - converged window:

weighting archetype	Balanced	Financial
Coal-heavy	0.791	2.079
Gas-dominant	0.548	1.364
Green-leader	0.067	0.172
Transitioner	0.560	0.313

1.8.11 RQ4.5 — Why agents go in compliant: per-cell root-cause attribution

The UDBC mix in §RQ3.4 / §RQ2.3 already tells us *how often* each combinatorial state (under-bid / debt / mixed / bank-covered / secondary-covered) occurs. What it does not isolate is the **proximate cause** of each non-compliant cell: did the agent bid too cheaply and lose the auction, did it ask for too few tonnes ($\text{bid_qty_mult} \times \text{estimate_need}$ short of realised need), did the

auction itself clear oversubscribed, was it carrying inherited debt forward, or did it submit a buy intent on the secondary that nobody filled?

This panel pulls every non-compliant cell (`episode`, `year`, `agent`) from the converged-tail window only and runs each through a prioritised waterfall:

1. **Inherited debt** — `alloc` `emiss` AND `cf_start` `shortfall`. The agent was auction-compliant *this year*; the deficit is fully explained by debt carried forward from earlier years.
2. **Underbid (price)** — `bid_price` < 99.5 % × `clearing_price`. The agent was priced out of the marginal allocation tranche.
3. **Undersized bid (quantity)** — (`bid_qty_mult` × `estimate_need`) falls short of the realised need by more than half of the shortfall. The agent asked for too few tonnes regardless of price.
4. **Oversubscribed auction** — `auction_total_demand` / `auction_volume` > 1.05. Even bidding at clearing, the agent was rationed by total demand exceeding supply.
5. **Secondary illiquidity** — agent submitted a positive buy intent (`sec_qty_action` > 0) but the realised buy fell more than 25 % of the shortfall short of intent.
6. **Mixed / cascade** — fallback bucket; no single cause dominates.

The taxonomy is intentionally simple and prioritised — each cell is assigned exactly one dominant cause. The downstream tables also report median diagnostic-feature values both per-cause and contrasting compliant vs non-compliant cells, so the *strength* of each signal is visible too.

```
[25]: # RQ4.5 - Per-cell root-cause attribution for non-compliant (agent, year)
# cells in the converged tail window. Cells silently no-op on logs lacking
# the v8.6.1+ year-log columns.
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

EPS_MT = 1e-6
CAUSE_ORDER = ['Inherited debt', 'Underbid (price)', 'Undersized bid (qty)',
               'Oversubscribed auction', 'Secondary illiquidity', 'Mixed /
               ↪ cascade']
CAUSE_COLORS = {
    'Inherited debt':      '#9467bd',
    'Underbid (price)':    '#d62728',
    'Undersized bid (qty)': '#ff7f0e',
    'Oversubscribed auction': '#1f77b4',
    'Secondary illiquidity': '#17becf',
    'Mixed / cascade':     '#7f7f7f',
}

def _classify_row(sf, cf_start, alloc, emiss, bid_p, clear_p,
                 undersized_qty, oversub, sec_unmet):
    if sf <= EPS_MT:
        return None
    cf = max(0.0, float(cf_start))
```

```

if alloc + EPS_MT >= emiss and cf >= sf - EPS_MT:
    return 'Inherited debt'
if bid_p > 0 and clear_p > 0 and bid_p < 0.995 * clear_p:
    return 'Underbid (price)'
if undersized_qty > 0.5 * sf:
    return 'Undersized bid (qty)'
if oversub > 1.05:
    return 'Oversubscribed auction'
if sec_unmet > 0.25 * sf:
    return 'Secondary illiquidity'
return 'Mixed / cascade'

nc_rows = []
for k, run in runs.items():
    yr = tail_year_rows(run)
    if yr.empty or 'year' not in yr.columns:
        continue
    n = len(yr)
    cp_arr = yr['clearing_price'].to_numpy(float) if 'clearing_price' in yr.columns else np.full(n, np.nan)
    atd = yr['auction_total_demand'].to_numpy(float) if 'auction_total_demand' in yr.columns else np.full(n, np.nan)
    avol = yr['auction_volume'].to_numpy(float) if 'auction_volume' in yr.columns else np.full(n, np.nan)
    ns_exec = yr['secondary_n_sellers_executed'].to_numpy(float) if 'secondary_n_sellers_executed' in yr.columns else np.full(n, np.nan)
    yr_arr = yr['year'].to_numpy(int)
    ep_arr = yr['episode'].to_numpy(int) if 'episode' in yr.columns else np.zeros(n, dtype=int)
    oversub_y = np.where(avol > 1e-9, atd / avol, np.nan)

    for i in range(N_AGENTS):
        cols_req = [f'shortfall_A{i+1}', f'alloc_A{i+1}', f'emissions_A{i+1}',
                    f'carry_forward_start_A{i+1}', f'bid_price_A{i+1}',
                    f'estimate_need_A{i+1}', f'bid_qty_mult_A{i+1}',
                    f'bank_start_A{i+1}', f'trade_qty_A{i+1}',
                    f'sec_qty_action_A{i+1}']
        if not all(c in yr.columns for c in cols_req):
            continue
        sf = yr[f'shortfall_A{i+1}'].to_numpy(float)
        nc_mask = sf > EPS_MT
        if not nc_mask.any():
            continue
        alloc = yr[f'alloc_A{i+1}'].to_numpy(float)
        emiss = yr[f'emissions_A{i+1}'].to_numpy(float)
        cf_st = yr[f'carry_forward_start_A{i+1}'].to_numpy(float)

```



```

bidp = yr[f'bid_price_A{i+1}'].to_numpy(float)
en = yr[f'estimate_need_A{i+1}'].to_numpy(float)
qm = yr[f'bid_qty_mult_A{i+1}'].to_numpy(float)
bank = yr[f'bank_start_A{i+1}'].to_numpy(float)
trq = yr[f'trade_qty_A{i+1}'].to_numpy(float)
sqa = yr[f'sec_qty_action_A{i+1}'].to_numpy(float)
cf_e = yr[f'carry_forward_end_A{i+1}'].to_numpy(float) if_
↳f'carry_forward_end_A{i+1}' in yr.columns else np.full(n, np.nan)
loan = yr[f'loan_outstanding_A{i+1}'].to_numpy(float) if_
↳f'loan_outstanding_A{i+1}' in yr.columns else np.full(n, np.nan)
tdr = yr[f'treasury_drawn_A{i+1}'].to_numpy(float) if_
↳f'treasury_drawn_A{i+1}' in yr.columns else np.full(n, np.nan)
cgap = yr[f'coverage_gap_A{i+1}'].to_numpy(float) if_
↳f'coverage_gap_A{i+1}' in yr.columns else np.full(n, np.nan)

# need_realised = emissions + inherited debt the agent had to cover.
need_realised = emiss + np.maximum(0.0, cf_st)
bid_qty_target = np.maximum(0.0, qm) * np.maximum(0.0, en)
undersized_qty = np.maximum(0.0, need_realised - bid_qty_target)
sec_intent_buy = np.maximum(0.0, sqa)
sec_bought = np.maximum(0.0, trq)
sec_unmet = np.maximum(0.0, sec_intent_buy - sec_bought)

for j in np.where(nc_mask)[0]:
    cp_j = float(cp_arr[j]) if np.isfinite(cp_arr[j]) else 0.0
    os_j = float(oversub_y[j]) if np.isfinite(oversub_y[j]) else 0.0
    cause = _classify_row(
        sf=float(sf[j]), cf_start=float(cf_st[j]),
        alloc=float(alloc[j]), emiss=float(emiss[j]),
        bid_p=float(bidp[j]), clear_p=cp_j,
        undersized_qty=float(undersized_qty[j]),
        oversub=os_j, sec_unmet=float(sec_unmet[j]),
    )
    nc_rows.append({
        'run': f'{k[1]}/s{k[2]}',
        'episode': int(ep_arr[j]),
        'year': int(yr_arr[j]),
        'agent': i + 1,
        'archetype': _agent_archetype(i),
        'weighting': _agent_weighting(i),
        'shortfall_Mt': float(sf[j]),
        'inherited_debt_Mt': float(cf_st[j]),
        'alloc_Mt': float(alloc[j]),
        'emissions_Mt': float(emiss[j]),
        'bid_price': float(bidp[j]),
        'clearing_price': cp_j,
        'bid_qty_mult': float(qm[j]),
    })

```

```

        'estimate_need_Mt': float(en[j]),
        'undersized_qty_Mt': float(undersized_qty[j]),
        'oversubscription_ratio': os_j,
        'bank_start_Mt': float(bank[j]),
        'sec_buy_intent_Mt': float(sec_intent_buy[j]),
        'secondary_buy_unmet_Mt': float(sec_unmet[j]),
        'n_sellers_executed': float(ns_exec[j]) if np.
↳ isfinite(ns_exec[j]) else np.nan,
        'carry_forward_end_Mt': float(cf_e[j]),
        'loan_outstanding_Meur': float(loan[j]),
        'treasury_drawn_Meur': float(tdr[j]),
        'coverage_gap_Mt': float(cgap[j]),
        'cause': cause,
    })

incomp_df = pd.DataFrame(nc_rows)

if incomp_df.empty:
    print(['RQ4.5] No non-compliant (agent, year) cells in the converged window_
↳ - '
        'agents are fully compliant. Skipping attribution panels.'])
else:
    incomp_df['cause'] = pd.Categorical(incomp_df['cause'],
↳ categories=CAUSE_ORDER, ordered=True)

    # (a) Counts and tonne-weighted shares per cause.
    cause_counts = (incomp_df.groupby('cause', observed=True)
        .agg(n_cells=('shortfall_Mt', 'size'),
            sum_shortfall_Mt=('shortfall_Mt', 'sum'),
            median_shortfall_Mt=('shortfall_Mt',
↳ 'median'))
        .reindex(CAUSE_ORDER).fillna(0.0))
    total_cells = max(1, int(cause_counts['n_cells'].sum()))
    total_tons = max(1e-9, float(cause_counts['sum_shortfall_Mt'].sum()))
    cause_counts['cell_share_pct'] = 100.0 * cause_counts['n_cells'] /
↳ total_cells
    cause_counts['tonne_share_pct'] = 100.0 * cause_counts['sum_shortfall_Mt'] /
↳ total_tons

    print(f' [RQ4.5] Non-compliant (agent, year, episode) cells in converged_
↳ window: '
        f'{len(incomp_df):,}'))
    print(f' [RQ4.5] Σ shortfall across the window: '
        f'{incomp_df["shortfall_Mt"].sum():.2f} Mt')
    display(cause_counts.round(3))

```

```

fig, axes = plt.subplots(1, 2, figsize=(14, 4.6))
colors = [CAUSE_COLORS[c] for c in CAUSE_ORDER]
axes[0].bar(CAUSE_ORDER, cause_counts['cell_share_pct'].values,
color=colors)
axes[0].set(title='Share of non-compliant cells by dominant cause',
ylabel='% of (agent, year, episode) cells')
axes[0].tick_params(axis='x', rotation=25); axes[0].grid(alpha=0.3,
axis='y')

axes[1].bar(CAUSE_ORDER, cause_counts['tonne_share_pct'].values,
color=colors)
axes[1].set(title='Share of total shortfall tonnes by dominant cause',
ylabel='% of  $\Sigma$  shortfall (Mt)')
axes[1].tick_params(axis='x', rotation=25); axes[1].grid(alpha=0.3,
axis='y')
plt.suptitle('RQ4.5 - Why agents go non-compliant in the converged window',
y=1.03, fontsize=13)
plt.tight_layout(); plt.show()

# (b) Median diagnostic feature values per cause - what discriminates them.
diag_cols = ['shortfall_Mt', 'inherited_debt_Mt',
'bid_price', 'clearing_price', 'bid_qty_mult',
'undersized_qty_Mt', 'oversubscription_ratio',
'bank_start_Mt', 'sec_buy_intent_Mt',
'secondary_buy_unmet_Mt', 'loan_outstanding_Meur']
diag_cols = [c for c in diag_cols if c in incomp_df.columns]
cause_medians = (incomp_df.groupby('cause', observed=True)[diag_cols]
.median().reindex(CAUSE_ORDER))
print('[RQ4.5] Median feature value per dominant cause:')
display(cause_medians.round(3))

# (c) Per-archetype  $\times$  weighting cause split.
aw_pivot = (incomp_df.groupby(['archetype', 'weighting', 'cause'],
observed=True)
.size().unstack('cause', fill_value=0))
aw_share = aw_pivot.div(aw_pivot.sum(axis=1).replace(0, np.nan), axis=0).
fillna(0.0)
aw_share = aw_share.reindex(columns=[c for c in CAUSE_ORDER if c in
aw_share.columns])
if not aw_share.empty:
fig, ax = plt.subplots(figsize=(11, 4.2))
aw_share.plot(kind='bar', stacked=True, ax=ax,
color=[CAUSE_COLORS[c] for c in aw_share.columns])
ax.set(title='RQ4.5 - Cause mix by archetype  $\times$  weighting (share within
row)',
ylabel='Share of non-compliant cells', ylim=(0, 1.0))

```

```

ax.legend(fontsize=8, loc='center left', bbox_to_anchor=(1.02, 0.5))
ax.tick_params(axis='x', rotation=20); ax.grid(alpha=0.3, axis='y')
plt.tight_layout(); plt.show()

# (d) Compliant vs non-compliant feature contrast on the same tail window.
comp_pieces = []
for k, run in runs.items():
    yr = tail_year_rows(run)
    if yr.empty: continue
    for i in range(N_AGENTS):
        sf_col = f'shortfall_A{i+1}'
        if sf_col not in yr.columns: continue
        sf = yr[sf_col].to_numpy(float)
        cmask = sf <= EPS_MT
        if not cmask.any(): continue
        row = {}
        for stub, src_col in [
            ('inherited_debt_Mt', f'carry_forward_start_A{i+1}'),
            ('bank_start_Mt', f'bank_start_A{i+1}'),
            ('bid_qty_mult', f'bid_qty_mult_A{i+1}'),
            ('bid_price', f'bid_price_A{i+1}'),
            ('coverage_gap_Mt', f'coverage_gap_A{i+1}'),
        ]:
            row[stub] = (yr.loc[cmask, src_col].to_numpy(float)
                         if src_col in yr.columns else np.full(int(cmask.
→sum()), np.nan))
            row['clearing_price'] = (yr.loc[cmask, 'clearing_price'].
→to_numpy(float)
                                   if 'clearing_price' in yr.columns
                                   else np.full(int(cmask.sum()), np.nan))
            comp_pieces.append(pd.DataFrame(row))
    if comp_pieces:
        comp_df = pd.concat(comp_pieces, ignore_index=True)
        comp_df['bid_to_clearing'] = (comp_df['bid_price']
→ / comp_df['clearing_price'].replace(0,
→np.nan))
        nc_with = incomp_df.copy()
        nc_with['bid_to_clearing'] = (nc_with['bid_price']
→ / nc_with['clearing_price'].replace(0,
→np.nan))
        contrast_cols = ['inherited_debt_Mt', 'bank_start_Mt',
                         'bid_qty_mult', 'bid_to_clearing', 'coverage_gap_Mt']
        rows = []
        for c in contrast_cols:
            cv = comp_df[c].replace([np.inf, -np.inf], np.nan).dropna() if c in
→comp_df.columns else pd.Series(dtype=float)

```

```

nv = nc_with[c].replace([np.inf, -np.inf], np.nan).dropna() if c in
nc_with.columns else pd.Series(dtype=float)
rows.append({
    'feature': c,
    'compliant_median': float(cv.median()) if len(cv) else np.
nan,
    'non_compliant_median': float(nv.median()) if len(nv) else np.
nan,
    'delta_(nc - c)': (float(nv.median() - cv.median())
                        if len(cv) and len(nv) else np.nan),
})
print('RQ4.5 Compliant vs non-compliant cell medians (key_
discriminators):')
display(pd.DataFrame(rows).round(3))


```

[RQ4.5] Non-compliant (agent, year, episode) cells in converged window: 13,477

[RQ4.5] Σ shortfall across the window: 16,309.33 Mt

	n_cells	sum_shortfall_Mt	median_shortfall_Mt	
cell_share_pct	tonne_share_pct			
cause				
Inherited debt	1484.0	2259.572	1.663	
11.011	13.854			
Underbid (price)	828.0	1259.363	0.793	
6.144	7.722			
Undersized bid (qty)	10055.0	11611.829	1.105	
74.609	71.197			
Oversubscribed auction	1107.0	1177.462	0.706	
8.214	7.220			
Secondary illiquidity	0.0	0.000	0.000	
0.000	0.000			
Mixed / cascade	3.0	1.103	0.352	
0.022	0.007			


RQ4.5 — Why agents go non-compliant in the converged window



[RQ4.5] Median feature value per dominant cause:

		shortfall_Mt	inherited_debt_Mt	bid_price	
clearing_price	bid_qty_mult	undersized_qty_Mt	oversubscription_ratio		
bank_start_Mt	sec_buy_intent_Mt	secondary_buy_unmet_Mt			
cause					
Inherited debt		1.663	1.998	165.621	60.
↳751	0.500	1.715	1.543	0.000	↳
↳	3.0	3.0			
Underbid (price)		0.793	0.000	45.000	97.
↳609	2.000	0.000	1.776	0.139	↳
↳	3.0	3.0			
Undersized bid (qty)		1.105	0.075	136.312	45.
↳000	0.500	1.540	1.247	0.000	↳
↳	3.0	3.0			
Oversubscribed auction		0.706	0.000	45.000	45.
↳000	2.000	0.000	1.432	0.187	↳
↳	3.0	3.0			
Secondary illiquidity		NaN	NaN	NaN	↳
↳NaN	NaN	NaN	NaN	NaN	↳
↳	NaN	NaN			
Mixed / cascade		0.352	0.000	45.000	45.
↳000	0.957	0.000	1.032	0.000	↳
↳	0.0	0.0			

	loan_outstanding_Meur
cause	
Inherited debt	0.0
Underbid (price)	0.0
Undersized bid (qty)	0.0
Oversubscribed auction	0.0
Secondary illiquidity	NaN
Mixed / cascade	0.0



[RQ4.5] Compliant vs non-compliant cell medians (key discriminators):

	feature	compliant_median	non_compliant_median	delta_(nc - c)
0	inherited_debt_Mt	0.000	0.182	0.182
1	bank_start_Mt	1.945	0.000	-1.945
2	bid_qty_mult	1.823	0.500	-1.323
3	bid_to_clearing	1.540	2.249	0.709
4	coverage_gap_Mt	0.000	2.118	2.118

RQ4 summary. RQ4.1–4.3 give the run-level mechanics: did the agents close the coverage gap, did the auction stay diffuse, and did the secondary market actually clear the residual? RQ4.4 attributes the residual debt and emergency-loan usage to the (archetype $\times$ weighting) cells — the financial vs ESG split is the most informative cut here, since coal-heavy financial agents are the ones for whom the gap-penalty reward (driven by `effective_penalty_rate`) actively dominates the bid head.



1.9 What this notebook does not answer

- Cross-config sensitivity — needs sweep variants. *Out of scope here.*
- Causal attribution of regulatory mechanisms (e.g. counterfactual “what if MSR were off”) — needs sweep variants.
- Real-world EU-ETS empirical validation beyond the qualitative price-band check in RQ1.3.
- Strategy clustering, theory tests (LCOE crossover, archetype heterogeneity), and predictive models of clearing price — these are in the analysis plan (`planning/data_science_analysis_plan.md` §B1, B2, B4) and will get their own notebooks where they can be evaluated cleanly with cross-validation.

The default-config evidence base assembled above is the *baseline* that those follow-on analyses will perturb and test against.