

ets_marl - Data Science Analysis

May 5, 2026

1 ETS MARL — Data Science Analysis

Implements the analyses described in `planning/data_science_analysis_plan.md` (v3). Every section ties back to a specific thesis claim. The structure is:

§	Analysis	Claim	Methods
C1	Convergence	RQ1.a	Rolling-mean plateau + PELT change-point
C2	Cross-seed reproducibility	RQ1.b	Fan charts + cross-seed CV
S1	Strategy archetypes	RQ3.a	KMeans + hierarchical + PCA
S2	Reward-weight effect on strategy	RQ3.b	S1 pipeline applied to all-financial variant; centroid contrast
S3	Economic rationality of strategies	RQ3.c	LASSO logit + Random Forest + anchor-vs-realised price
R1	Regulatory shifts	RQ2.a	Cluster shift in shared PCA space + outcome boxplots (IQR)
R2	UDBC compliance pathways	RQ2.b	Multinomial logit + RF + SMOTE; transition heatmaps

Data ingestion. Mirrors `notebooks/ets_marl - Default RQ Analysis.ipynb` and `notebooks/ets_marl - Sweep Analysis.ipynb`: load one or more sweep specs, expand each variant's overrides over the base config, then walk `<output_dir>/<variant>/` for `training_log_*.csv` (per episode) and `year_log_*.csv` (per year). All runs are keyed by `(sweep, variant, seed)`.

Robustness. Every section guards against missing data — sections that need a non-default variant (S2 cross-variant, R1, R2 cross-variant) skip gracefully when only default seeds are loaded. This means the notebook runs end-to-end on the smoke dataset (`configs/sweeps/smoke_ds_analysis.yaml`) and on the full sweep without edits.

Methodology. Standardisation before any distance-based method (KMeans, PCA, LASSO). Train/val splits or cross-validation for any predictive model. Multi-method robustness on every

headline claim — see the table in the planning doc §7.

```
[1]: # Install notebook dependencies if needed.
# Set INSTALL_NOTEBOOK_DEPS = True (or run the cell with required packages
↳missing)
# to pip-install everything in notebooks/requirements-notebooks.txt using the
# CURRENT kernel's interpreter. This is more reliable than `%pip install -r
↳{path}`
# (IPython magics don't expand `{var}` at runtime) and works on Azure / Colab /
↳local.
import subprocess, sys, importlib, re
from pathlib import Path

INSTALL_NOTEBOOK_DEPS = False # flip to True to force-reinstall

# Map PyPI distribution name -> import module name when they differ.
_DIST_TO_IMPORT = {
    "pyyaml": "yaml",
    "scikit-learn": "sklearn",
    "imbalanced-learn": "imblearn",
    "opencv-python": "cv2",
    "pillow": "PIL",
    "beautifulsoup4": "bs4",
    "ipython": "IPython",
}

def _find_requirements():
    req_name = "requirements-notebooks.txt"
    here = Path.cwd()
    candidates = [here / req_name, here / "notebooks" / req_name]
    candidates.extend(p / "notebooks" / req_name for p in here.parents)
    return next((p for p in candidates if p.exists()), None)

def _parse_requirements(req_path):
    """Extract distribution names from a requirements file (skip comments,
↳options, env markers)."""
    dists = []
    for raw in req_path.read_text(encoding="utf-8").splitlines():
        line = raw.strip()
        if not line or line.startswith("#") or line.startswith("-"):
            continue
        # Strip env markers and extras: "pkg[extra]>=1.0; python_version>='3.9'"
        line = line.split(";", 1)[0].strip()
        m = re.match(r"^(?![A-Za-z0-9_.\-]+)", line)
        if m:
            dists.append(m.group(1).lower())
    return dists
```

```

def _missing_required(dists):
    missing = []
    for dist in dists:
        mod = _DIST_TO_IMPORT.get(dist, dist.replace("-", "_"))
        try:
            importlib.import_module(mod)
        except ImportError:
            missing.append(dist)
    return missing

req_path = _find_requirements()
if req_path is None:
    raise FileNotFoundError("Could not find notebooks/requirements-notebooks.
↪txt")

dists = _parse_requirements(req_path)
missing = _missing_required(dists)

if INSTALL_NOTEBOOK_DEPS or missing:
    if missing:
        print(f"Missing packages: {missing}")
        print(f"Installing notebook dependencies from: {req_path}")
        subprocess.check_call([sys.executable, "-m", "pip", "install", "-r",
↪str(req_path)])
    else:
        print(f"All {len(dists)} notebook dependencies present; skipping install.")
        print("(Set INSTALL_NOTEBOOK_DEPS = True to force a reinstall.)")

```

All 14 notebook dependencies present; skipping install.
(Set INSTALL_NOTEBOOK_DEPS = True to force a reinstall.)

1.1 1. Setup — Locate Project Root

```

[2]: import os, sys
from pathlib import Path

def _is_project_root(path: str) -> bool:
    return os.path.isdir(os.path.join(path, "src")) and os.path.isdir(os.path.
↪join(path, "configs"))

_candidates = ["..", ".", os.path.join("..", ".."),
               os.path.join("Thesis-Energy-Auction", "ets_marl_happo_current"),
↪"Thesis-Energy-Auction"]
_found = _is_project_root(".")
if not _found:
    for _c in _candidates:

```

```

        if _is_project_root(_c):
            os.chdir(_c); _found = True; break
if not _found:
    raise RuntimeError(f"Could not locate project root from {os.getcwd()}.")

PROJECT_ROOT = Path(os.getcwd()).resolve()
if str(PROJECT_ROOT) not in sys.path:
    sys.path.insert(0, str(PROJECT_ROOT))
print("Project root:", PROJECT_ROOT)

```

Project root: C:\Users\danie\Documents\Python_Scripts\Master Thesis\Thesis-Energy-Auction

1.2 2. Imports

Standard scientific stack plus a few sklearn / specialised modules used by the analyses below.

```

[3]: import re
import yaml
import warnings
import itertools
from pathlib import Path
from collections import Counter, defaultdict

import numpy as np
import pandas as pd
from src.utils.run_data import load_run_csv, glob_run_logs # parquet-cached
↳ log loader
import matplotlib.pyplot as plt
import seaborn as sns

# ML / DS toolbox
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
from sklearn.cluster import KMeans, AgglomerativeClustering
from sklearn.metrics import silhouette_score, classification_report
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.inspection import partial_dependence
from sklearn.model_selection import StratifiedKFold, cross_val_score,
↳ train_test_split
from sklearn.pipeline import Pipeline
from scipy.cluster.hierarchy import linkage, dendrogram, fcluster

# Specialised - change-point detection (justified in methodology) and
# minority-class oversampling for R2.
import ruptures as rpt

```

```

try:
    from imblearn.over_sampling import SMOTE
    _HAS_SMOTE = True
except ImportError:
    _HAS_SMOTE = False

# Project utilities - sweep loader and price anchor.
from src.utils.sweep import load_sweep_spec, expand_dotted_overrides, u
↳ deep_merge
from src.utils.price_anchor import compute_fundamental_anchor

sns.set_style('whitegrid')
pd.set_option('display.max_columns', 200)
pd.set_option('display.width', 200)
warnings.filterwarnings('ignore', category=FutureWarning)
warnings.filterwarnings('ignore', category=UserWarning, module='sklearn')

RNG = np.random.default_rng(0)
print('Imports OK. SMOTE available:', _HAS_SMOTE)

```

Imports OK. SMOTE available: True

1.3 3. Sweep Specs to Load

List one or more sweep specs (paths relative to the project root). Each spec's `output_dir` may be overridden if results were moved. Multiple specs can be combined — runs from every spec are merged into a single (sweep, variant, seed)-keyed dict.

To validate the notebook end-to-end on a tiny dataset, run `python scripts/sweep.py --spec configs/sweeps/smoke_ds_analysis.yaml` first, then keep `SWEEP_SPECS` pointed at it.

To analyse a real run, swap in the production sweep spec(s) (e.g. `configs/sweeps/default_seeds.yaml`).

```

[4]: # Sweep specs to load. Add or comment-out entries as the experiments roll in.
#
#   default_seeds           - RQ1 (4 seeds at default config)
#   thesis_experiments      - RQ2 (LRF / MSR / inflation) and RQ3.b (all_fin /
#                               balanced); 16 (variant, seed) jobs total. Uncomment
#                               once the sweep finishes.
SWEEP_SPECS = [
    PROJECT_ROOT / 'configs' / 'sweeps' / 'default_seeds.yaml',
    PROJECT_ROOT / 'configs' / 'sweeps' / 'thesis_experiments.yaml',
    # PROJECT_ROOT / 'configs' / 'sweeps' / 'smoke_ds_analysis.yaml',
]
# `default_seeds.yaml` is loaded with all CSVs in a single `default/` folder
# (no per-variant subfolder), so we override its output_dir to point at that
# flat layout. The thesis sweep follows the standard one-folder-per-variant
# layout and needs no override (pass `None`).

```

```

OUTPUT_DIR_OVERRIDES = [
    PROJECT_ROOT / 'results' / 'sweeps' / 'default',
    None,      # thesis_experiments
    # None,    # smoke_ds_analysis
]
assert len(SWEEP_SPECS) == len(OUTPUT_DIR_OVERRIDES)

# -----
# Memory-budget knobs.
#
# The notebook NEVER holds more than one run's raw DataFrames in RAM at a
# time. A single streaming pass (see "Streaming Pre-compute" below) loads
# each run, distils the small artefacts that downstream sections need
# (decimated traces, per-(run, agent) feature rows, per-(run, year)
# clearing prices, UDBC labels, ...), and drops the run before moving on.
#
# ``EXPERIMENT_FILTER`` lets you subset which (sweep, variant, seed)
# triples participate without editing SWEEP_SPECS. Use ``None`` or
# ``['*']`` to pass everything; otherwise supply fnmatch-style patterns.
# This is the knob to reach for when even the streaming pass produces
# more derived artefacts than your machine has RAM for - e.g. drop
# ``high_inf`` while you focus on RQ3.b, then swap back later.
# -----
EXPERIMENT_FILTER = {
    'sweeps': None,      # e.g. ['default_seeds', 'thesis_experiments']
    'variants': None,    # e.g. ['default', 'all_fin*', 'balanced']
    'seeds': None,       # e.g. [42, 123] or range(0, 4)
    'max_runs': None,    # hard cap; None = unlimited
}

# Convergence-window definition. ``TAIL_FRAC`` of the longest run is used
# for steady-state metrics; this is overridden per-run by the C1 detector
# when both PELT and the rolling-plateau detector agree.
TAIL_FRAC = 0.20

# Pre-compute streaming budgets. Lower these if memory is tight.
# PRECOMPUTE_TAIL_EPISODES - episode-log rows kept per run during the
# streaming pass. None = the full episode log is read; an int caps
# the read to the last N episodes (parquet row-filter pushdown,
# cheap). The detectors and plot traces decimate from whatever is
# read, so a value below the run length still gives faithful
# steady-state metrics but loses the convergence-warmup view in
# the SC1 plot. A safe default is ``None`` for runs <= 50k episodes
# and 50_000 for longer runs.
PRECOMPUTE_TAIL_EPISODES = None

```

```
# UDBC fit cap (subsamped from the converged-window panel before SR2
# logit/RF; see cell SR2). Set lower if SR2 alone is the bottleneck.
R2_FIT_CAP = 200_000
```

1.4 4. Discover Runs (Registry Only)

Build a lightweight registry of every (sweep, variant, seed) triple that the sweep specs above promise. **No raw DataFrames are loaded here** - just paths, the effective per-variant config, and a parquet metadata peek for episode counts. The registry is cheap (kilobytes) even for hundreds of runs, so it fits in RAM regardless of dataset size.

The streaming pre-compute cell further down does the heavy lifting one run at a time. Apply EXPERIMENT_FILTER (see §3) to subset which runs participate.

```
[5]: import fnmatch
import gc
from src.utils.run_data import (
    log_exists, default_ep_columns, default_yr_columns,
    cache_path_for, load_run_csv,
)
from tqdm.auto import tqdm

# Project both logs onto the analysis-superset columns at load time. The
# year_log carries 270+ diagnostic columns; pulling only the ~80 actually
# used by this notebook cuts in-RAM size by ~5-10x. Unknown columns are
# silently dropped by load_run_csv.
EP_COLUMNS = default_ep_columns()
YR_COLUMNS = default_yr_columns()

def _resolve_run_files(variant_name, seed, variant_dir, output_dir):
    """Return (training_log, year_log) paths or (None, None). Tagged then
    ↪untagged."""
    search_dirs = [d for d in (variant_dir, output_dir) if d is not None]
    training_names = [f"training_log_{variant_name}_s{seed}.csv",
    ↪f"training_log_s{seed}.csv"]
    year_names = [f"year_log_{variant_name}_s{seed}.csv",
    ↪f"year_log_s{seed}.csv"]
    tlog = ylog = None
    for base in search_dirs:
        if tlog is None:
            for n in training_names:
                p = base / n
                if log_exists(p):
                    tlog = p; break
    if ylog is None:
```

```

        for n in year_names:
            p = base / n
            if log_exists(p):
                ylog = p; break
    return tlog, ylog

def _peek_n_episodes(csv_path):
    """Cheap row-count via parquet metadata; falls back to None on miss."""
    if csv_path is None:
        return None
    try:
        import pyarrow.parquet as pq
        pq_path = cache_path_for(str(csv_path))
        if Path(pq_path).exists():
            return int(pq.ParquetFile(pq_path).metadata.num_rows)
    except Exception:
        pass
    return None

def _filter_match(value, patterns):
    """fnmatch-style match against a list of patterns; ``None``/'`['*']`` = 
    ↪pass."""
    if patterns is None:
        return True
    if isinstance(patterns, str):
        patterns = [patterns]
    if any(p == '*' for p in patterns):
        return True
    return any(fnmatch.fnmatch(str(value), p) for p in patterns)

def _seed_match(seed, allowed):
    if allowed is None:
        return True
    return int(seed) in {int(s) for s in allowed}

def _build_registry(spec_path, output_override, exp_filter):
    spec = load_sweep_spec(str(spec_path))
    base_cfg_path = (PROJECT_ROOT / spec['base_config']).resolve()
    with open(base_cfg_path) as f:
        base_cfg = yaml.safe_load(f)
    output_dir = Path(output_override) if output_override else (PROJECT_ROOT / 
    ↪spec['output_dir']).resolve()
    sweep_name = Path(spec_path).stem

```

```

if not _filter_match(sweep_name, exp_filter.get('sweeps')):
    return {}

entries = {}
for v in spec['variants']:
    if not _filter_match(v['name'], exp_filter.get('variants')):
        continue
    seeds = v.get('seeds') or spec.get('seeds') or []
    overrides = v.get('overrides') or {}
    eff_cfg = deep_merge(base_cfg, expand_dotted_overrides(overrides))
    variant_dir = output_dir / v['name']
    for s in seeds:
        if not _seed_match(s, exp_filter.get('seeds')):
            continue
        tlog, ylog = _resolve_run_files(v['name'], int(s), variant_dir,
↪output_dir)
        if tlog is None:
            continue
        entries[(sweep_name, v['name'], int(s))] = {
            'sweep': sweep_name,
            'variant': v['name'],
            'seed': int(s),
            'has_overrides': bool(overrides),
            'overrides': overrides,
            'config': eff_cfg,
            'training_log': str(tlog),
            'year_log': str(ylog) if ylog else None,
            'n_episodes': _peek_n_episodes(tlog), # cheap parquet↪
↪metadata read
        }
    return entries

# Build registry across all SWEEP_SPECS (and apply EXPERIMENT_FILTER).
runs = {}
for spec_path, override in zip(SWEEP_SPECS, OUTPUT_DIR_OVERRIDES):
    runs.update(_build_registry(spec_path, override, EXPERIMENT_FILTER))

# Apply max_runs cap (sorted for determinism).
_max = EXPERIMENT_FILTER.get('max_runs')
if _max is not None and len(runs) > _max:
    keep_keys = sorted(runs.keys())[:int(_max)]
    runs = {k: runs[k] for k in keep_keys}

if not runs:
    raise RuntimeError(

```

```

        'No runs in registry after filtering. Check SWEEP_SPECS / '
        'OUTPUT_DIR_OVERRIDES / EXPERIMENT_FILTER.'
    )

print(f'Registry built: {len(runs)} run(s) (no raw DataFrames loaded yet).')
for k, r in runs.items():
    flag = '[overrides]' if r['has_overrides'] else '[default]'
    n_ep = r['n_episodes'] if r['n_episodes'] is not None else '?'
    print(f"    {k}    n_ep={n_ep}    {flag}")

# -----
# Lazy per-run loader. Use this for ad-hoc analysis on a single run -
# the streaming pre-compute below covers every section in this
# notebook, so reach for ``load_run_data`` only when poking outside
# the pre-built artefacts.
# -----
def load_run_data(key, *, ep_columns=EP_COLUMNS, yr_columns=YR_COLUMNS,
                  tail_episodes=None, with_yr=True):
    """Load one run's training_log + year_log on demand.

    The caller owns the returned DataFrames - ``del`` and ``gc.collect()``
    when done so peak RAM stays bounded. Pass ``tail_episodes=N`` to keep
    only the last N episodes (parquet row-filter pushdown).
    """
    entry = runs[key]
    ep_df = load_run_csv(entry['training_log'], columns=ep_columns,
                        tail_episodes=tail_episodes)
    yr_df = pd.DataFrame()
    if with_yr and entry['year_log']:
        yr_df = load_run_csv(entry['year_log'], columns=yr_columns,
                            tail_episodes=tail_episodes)
    return {**entry, 'ep_df': ep_df, 'yr_df': yr_df if yr_df is not None else
    ↪pd.DataFrame()}

def iter_run_data(filter_keys=None, *, ep_columns=EP_COLUMNS,
    ↪yr_columns=YR_COLUMNS,
                  tail_episodes=None, with_yr=True, desc='loading'):
    """Generator: yields ``(key, run_dict)`` one run at a time, dropping
    each before yielding the next. Use in ``for k, run in iter_run_data():``
    streaming loops to keep peak RAM at one run.
    """
    keys = list(filter_keys) if filter_keys is not None else list(runs.keys())
    pbar = tqdm(keys, desc=desc, unit='run', leave=False)
    for k in pbar:
        run = load_run_data(k, ep_columns=ep_columns, yr_columns=yr_columns,

```

```

tail_episodes=tail_episodes, with_yr=with_yr)

yield k, run
del run
gc.collect()

```

Registry built: 20 run(s) (no raw DataFrames loaded yet).

```

('default_seeds', 'default', 1729)  n_ep=120000  [default]
('default_seeds', 'default', 8191)  n_ep=120000  [default]
('default_seeds', 'default', 6561)  n_ep=120000  [default]
('default_seeds', 'default', 5041)  n_ep=120000  [default]
('thesis_experiments', 'lrf_low', 1729)  n_ep=120000  [overrides]
('thesis_experiments', 'lrf_low', 8191)  n_ep=120000  [overrides]
('thesis_experiments', 'lrf_low', 6561)  n_ep=120000  [overrides]
('thesis_experiments', 'lrf_high', 1729)  n_ep=120000  [overrides]
('thesis_experiments', 'lrf_high', 8191)  n_ep=120000  [overrides]
('thesis_experiments', 'lrf_high', 6561)  n_ep=120000  [overrides]
('thesis_experiments', 'no_msr', 1729)  n_ep=120000  [overrides]
('thesis_experiments', 'no_msr', 8191)  n_ep=120000  [overrides]
('thesis_experiments', 'no_msr', 6561)  n_ep=120000  [overrides]
('thesis_experiments', 'all_fin', 1729)  n_ep=120000  [overrides]
('thesis_experiments', 'all_fin', 8191)  n_ep=120000  [overrides]
('thesis_experiments', 'all_fin', 6561)  n_ep=120000  [overrides]
('thesis_experiments', 'balanced', 1729)  n_ep=120000  [overrides]
('thesis_experiments', 'balanced', 8191)  n_ep=120000  [overrides]
('thesis_experiments', 'balanced', 6561)  n_ep=120000  [overrides]
('thesis_experiments', 'high_inf', 1729)  n_ep=120000  [overrides]

```

1.5 5. Shared Helpers

These wrap the column-naming convention used throughout the trainer (`<metric>_A<i>` for agent-level columns, 1-indexed) and provide a few analysis-window primitives reused across sections.

```

[6]: # Number of learning agents and timeline length, taken from the first
# registry entry's config (no DataFrames needed).
_first_run = next(iter(runs.values()))
N_AGENTS = int(_first_run['config']['companies']['n_agents'])
N_YEARS = int(_first_run['config']['simulation']['n_years'])
print(f'N_AGENTS={N_AGENTS}  N_YEARS={N_YEARS}')

# Plot/detector decimation budgets. Production runs are 100k+ episodes; PELT
# on raw episodes takes minutes per series and matplotlib chokes on 100k-point
# scatters. Decimate to a fixed cap before any O(n) detector or plot. Mean-shift
# detection is unaffected by uniform decimation; the plateau cutoff is mapped
# back to the original episode index.
MAX_DETECTOR_POINTS = 5000
MAX_PLOT_POINTS = 4000

```

```

def decimate_series(arr, max_points):
    """Uniform-step decimation. Returns (kept_values, original_indices).

    ``arr`` may be a numpy array, list, or pandas Series. If ``len(arr) <=
    max_points`` the original is returned unchanged with ``np.arange(len)``.
    """
    a = np.asarray(arr)
    n = len(a)
    if n <= max_points or max_points <= 0:
        return a, np.arange(n)
    step = max(1, n // max_points)
    idx = np.arange(0, n, step)
    return a[idx], idx

def map_decimated_index(idx_in_decimated, original_indices):
    """Map a position in the decimated array back to the original index."""
    if idx_in_decimated is None:
        return None
    idx_in_decimated = int(idx_in_decimated)
    if idx_in_decimated < 0 or idx_in_decimated >= len(original_indices):
        return None
    return int(original_indices[idx_in_decimated])

def agent_cols(df, prefix, max_idx=None):
    """Return ``[<prefix>_A1, ..., <prefix>_AN]`` for columns present in ``df``.
    ↪ """
    if max_idx is None: max_idx = N_AGENTS
    rx = re.compile(rf'^({re.escape(prefix)}_A(\d+)$)')
    out = []
    for c in df.columns:
        m = rx.match(c)
        if m and int(m.group(1)) <= max_idx:
            out.append((int(m.group(1)), c))
    return [c for _, c in sorted(out)]

def archetype_of(agent_idx_0based: int, config: dict) -> str:
    """Designed archetype label from initial mix (% green = onshore + offshore,
    ↪ + solar)."""
    mix = config['companies']['initial_mix'][agent_idx_0based]
    green = float(mix[2] + mix[3] + mix[4])
    if green < 0.50: return 'coal-leaning'
    if green < 0.60: return 'balanced'
    if green < 0.70: return 'transitioner'

```

```

return 'green-leader'

def weighting_of(agent_idx_0based: int, config: dict) -> str:
    w = config['companies']['reward_weights'][agent_idx_0based]
    if float(w[1]) <= 1e-6: return 'pure_financial'
    return 'balanced_esg'

def n_episodes_of(key):
    """Per-run episode count from the registry. Falls back to 0 if unknown."""
    return int(runs[key].get('n_episodes') or 0)

def converged_window(ep_df, tail_frac=TAIL_FRAC, override_start=None):
    """Return the steady-state slice of an episode log."""
    n = len(ep_df)
    if override_start is not None and 0 <= override_start < n:
        return ep_df.iloc[override_start:].copy()
    cut = max(1, int(np.floor(n * (1.0 - tail_frac))))
    return ep_df.iloc[cut:].copy()

def run_key(run):
    return (run['sweep'], run['variant'], run['seed'])

# Convenience: split runs by default vs override. Both are registry views
# (cheap dict-of-dicts; no DataFrames).
DEFAULT_RUNS = {k: v for k, v in runs.items() if not v['has_overrides']}
OVERRIDE_RUNS = {k: v for k, v in runs.items() if v['has_overrides']}
print(f'default-config runs: {len(DEFAULT_RUNS)}    override-config runs: {len(OVERRIDE_RUNS)}')
print('variants present:', sorted({k[1] for k in runs}))

```

```

N_AGENTS=8  N_YEARS=12
default-config runs: 4    override-config runs: 16
variants present: ['all_fin', 'balanced', 'default', 'high_inf', 'lrf_high',
'lrf_low', 'no_msr']

```

1.6 §5b - Streaming Pre-compute

Single pass over the registry: load each run, extract every artefact that downstream sections consume, drop the run, repeat. Peak RAM is bounded by **one run's DataFrames at a time** regardless of how many runs are in EXPERIMENT_FILTER.

Re-run this cell whenever you change EXPERIMENT_FILTER, SWEEP_SPECS, or

PRECOMPUTE_TAIL_EPISODES. All §C/§S/§R/§6 cells below read from the artefacts produced here.

Artefacts populated:

Artefact	Type	Used by
conv_starts[k]	int per run	§C1, every steady-state metric
conv_results	DataFrame	§C1 detector table
plot_traces[k]	dict of decimated arrays	§C1 trajectory plot
headline_traces_default[name]	list of decimated seed arrays	§C2 fan chart
headline_tail_means[k][name]	float	§C2 CV table, §6 export
feat_df	DataFrame (per run, agent)	§S1, §S2, §R1.a
green_means_per_run[k]	dict	§S2 outcome compare
clearing_per_year_df	DataFrame (per run, year)	§S3.a, §R1.c
pay_df	DataFrame (per run, agent, year)	§S3 H1/H2
udbc_df	DataFrame (per run, agent, year)	§R2
out_df	DataFrame (per run)	§R1.b
export_rows	list of dicts	§6 headline export

```
[7]: import time
import ruptures as rpt

# --- Detector primitives (operate on decimated arrays only) -----
def rolling_plateau(series, window=None, tol_rel=0.02, hold_windows=3):
    """First episode where the rolling mean stays within ``tol_rel`` of the
    final-window mean for ``hold_windows`` consecutive windows.
    """
    n = len(series)
    if n < 20: return None
    if window is None:
        window = max(20, n // 20)
    s = pd.Series(series).rolling(window, min_periods=window).mean().values
    final_mean = np.nanmean(s[-window:])
    if not np.isfinite(final_mean) or abs(final_mean) < 1e-9:
        band = max(1.0, abs(final_mean) * tol_rel + 1.0)
    else:
        band = abs(final_mean) * tol_rel
    in_band = np.abs(s - final_mean) <= band
    streak = 0
    for i, ok in enumerate(in_band):
        if not np.isfinite(s[i]):
            streak = 0; continue
        streak = streak + 1 if ok else 0
        if streak >= hold_windows * window:
            return max(0, i - hold_windows * window + 1)
```

```

return None

def pelt_change(series, pen=None):
    """PELT change-point on an L2 cost; returns the last change-point or None.

    Pass an already-decimated input - on 5k l2 points this is sub-second.
    """
    n = len(series)
    if n < 20: return None
    x = np.asarray(series, dtype=float)
    x = x[np.isfinite(x)]
    if len(x) < 20: return None
    if pen is None:
        pen = 3.0 * np.log(len(x)) * np.var(x)
    algo = rpt.Pelt(model='l2', min_size=max(5, len(x) // 50)).fit(x)
    cps = algo.predict(pen=pen)
    cps = [c for c in cps if c < len(x)]
    return cps[-1] if cps else None

def _detect_on_decimated(series, max_points=MAX_DETECTOR_POINTS):
    if len(series) == 0:
        return None, None
    dec, orig_idx = decimate_series(series, max_points)
    plat = rolling_plateau(dec)
    pelt = pelt_change(dec)
    return map_decimated_index(plat, orig_idx), map_decimated_index(pelt,
↪orig_idx)

# --- Feature column lists (used in feat_df rows) -----
EPISODE_FEATURES = ['bid_price', 'avg_bid_mult', 'inv_onshore_share',
                    'inv_offshore_share', 'inv_solar_share', 'green_frac',
                    'invest_cost', 'penalty',
                    'peak_loan_outstanding', 'peak_carry_forward',
                    'final_treasury_reserve']
YEAR_FEATURES     = ['sec_buy_vol', 'sec_sell_vol',
                    'coverage_gap', 'effective_penalty_rate',
                    'carry_forward_end', 'treasury_reserve',
                    'loan_outstanding']
UDBC_BUCKETS      = ['U', 'D', 'M', 'B', 'C']

HEADLINE_SPEC = [
    ('system_reward',      'sum', 'reward',      None),
    ('quality_score',      'sum',  None,          'quality_score'),

```

```

        ('clearing_price_last', 'sum', None,
↪ 'clearing_price_last'),
        ('mean_green_frac', 'mean', 'green_frac', None),
        ('mean_penalty', 'mean', 'penalty', None),
        ('mean_invest_cost', 'mean', 'invest_cost', None),
    ]

def _system_metric(df, agg, prefix=None, col=None):
    if prefix is not None:
        cols = agent_cols(df, prefix)
        if not cols: return pd.Series(dtype=float)
        x = df[cols]
        return x.sum(axis=1) if agg == 'sum' else x.mean(axis=1)
    if col is not None and col in df.columns:
        return df[col]
    return pd.Series(dtype=float)

def _udbc_panel_for_run(yrc, cfg, k):
    """Vectorised UDBC label assignment for one run's converged year_log slice.
    ↪ """
    n = len(yrc)
    if n == 0:
        return pd.DataFrame()
    pieces = []
    cap_arr = yrc['cap'].to_numpy(dtype=float) if 'cap' in yrc.columns else np.
↪ full(n, np.nan)
    cp_arr = yrc['clearing_price'].to_numpy(dtype=float) if 'clearing_price'
↪ in yrc.columns else np.full(n, np.nan)
    yr_arr = yrc['year'].to_numpy(dtype=int)
    for i in range(N_AGENTS):
        sf = yrc.get(f'shortfall_A{i+1}', pd.Series(np.zeros(n))).
↪ to_numpy(dtype=float)
        bank = yrc.get(f'bank_start_A{i+1}', pd.Series(np.zeros(n))).
↪ to_numpy(dtype=float)
        alloc = yrc.get(f'alloc_A{i+1}', pd.Series(np.zeros(n))).
↪ to_numpy(dtype=float)
        emiss = yrc.get(f'emissions_A{i+1}', pd.Series(np.zeros(n))).
↪ to_numpy(dtype=float)
        sec = yrc.get(f'secondary_net_A{i+1}', pd.Series(np.zeros(n))).
↪ to_numpy(dtype=float)
        non_comp = sf > 1e-6
        cf = np.maximum(0.0, bank)
        under = alloc < emiss
        has_cf = cf > 1e-6

```

```

labels = np.where(sec > 1e-6, 'B', 'C').astype(object)
u_mask = non_comp & under & ~has_cf
d_mask = non_comp & ~under & has_cf
m_mask = non_comp & ~u_mask & ~d_mask
labels = np.where(u_mask, 'U', labels)
labels = np.where(d_mask, 'D', labels)
labels = np.where(m_mask, 'M', labels)
pieces.append(pd.DataFrame({
    'sweep':      k[0], 'variant':  k[1], 'seed':      k[2],
    'agent':      i + 1,
    'year':       yr_arr,
    'archetype':  archetype_of(i, cfg),
    'weighting':  weighting_of(i, cfg),
    'cap':        cap_arr,
    'clearing_price': cp_arr,
    'bank_start': bank,
    'label':      labels,
}))
return pd.concat(pieces, ignore_index=True) if pieces else pd.DataFrame()

# --- Artefact containers -----
conv_starts          = {}
conv_results         = []
plot_traces          = {}
headline_traces_default = {name: [] for name, *_ in HEADLINE_SPEC}
headline_tail_means  = {}
green_means_per_run  = {}
n_episodes_per_run   = {}
feat_rows            = []
clearing_per_year_rows = []
pay_pieces           = []
udbc_pieces          = []
out_rows             = []
export_rows          = []

# Columns we actually need from the year_log during pre-compute. The
# canonical default_yr_columns superset is fine, but we keep it slim
# here to avoid the year_log's per-run RAM spike.
_PRECOMPUTE_YR_COLS = list(set(YR_COLUMNS) | {'year', 'episode', 'cap', 'clearing_price'}) | {
    f'{p}_A{i+1}' for p in ['shortfall', 'bank_start', 'alloc', 'emissions',
                           'secondary_net', 'invest_cost']
    for i in range(N_AGENTS)
}

```

```

_t0 = time.time()
print(f'Streaming pre-compute over {len(runs)} run(s) '
      f'(tail_episodes={PRECOMPUTE_TAIL_EPISODES}, peak RAM = 1 run)...')

for k in tqdm(list(runs.keys()), desc='pre-compute', unit='run', leave=True):
    entry = runs[k]
    ep_df = load_run_csv(entry['training_log'], columns=EP_COLUMNS,
                        tail_episodes=PRECOMPUTE_TAIL_EPISODES)
    yr_df = pd.DataFrame()
    if entry['year_log']:
        _yr_loaded = load_run_csv(entry['year_log'],
    ↪columns=_PRECOMPUTE_YR_COLS,
                                tail_episodes=PRECOMPUTE_TAIL_EPISODES)
        if _yr_loaded is not None:
            yr_df = _yr_loaded

    n_ep = len(ep_df)
    n_episodes_per_run[k] = n_ep
    # If the registry's parquet-metadata peek missed earlier, refresh now.
    runs[k]['n_episodes'] = n_ep

    # ---- SC1 detectors -----
    rew_cols = agent_cols(ep_df, 'reward')
    rew = ep_df[rew_cols].sum(axis=1).values if rew_cols else np.array([])
    px = ep_df['clearing_price_last'].values if 'clearing_price_last' in ep_df.
    ↪columns else np.array([])
    plat_r, pelt_r = _detect_on_decimated(rew)
    plat_p, pelt_p = _detect_on_decimated(px)
    candidates = [c for c in (plat_r, plat_p, pelt_r, pelt_p) if c is not None]
    cutoff = max(candidates) if candidates else max(1, n_ep // 2)
    conv_starts[k] = cutoff
    conv_results.append({
        'run': f'{k[1]}/s{k[2]}', 'n_ep': n_ep,
        'plateau_reward': plat_r, 'plateau_price': plat_p,
        'pelt_reward': pelt_r, 'pelt_price': pelt_p,
        'convergence_cutoff': cutoff,
        'tail_frac': round(1.0 - cutoff / max(1, n_ep), 3),
    })

    # ---- SC1 plot traces (decimated raw + rolling mean) -----
    def _trace(y):
        if len(y) == 0:
            return None
        y_dec, x_dec = decimate_series(y, MAX_PLOT_POINTS)
        smooth = pd.Series(y).rolling(max(5, len(y) // 30), min_periods=1).
    ↪mean()
        sm_dec, sx_dec = decimate_series(smooth.values, MAX_PLOT_POINTS)

```

```

        return {'y_dec': y_dec, 'x_dec': x_dec, 'sm_dec': sm_dec, 'sx_dec':  

        ↪sx_dec}  

        plot_traces[k] = {'reward': _trace(rew), 'price': _trace(px), 'cutoff':  

        ↪cutoff}  

        # ---- SC2 headline traces & tail means -----  

        tail_means_k = {}  

        for name, agg, prefix, col in HEADLINE_SPEC:  

            s = _system_metric(ep_df, agg, prefix=prefix, col=col)  

            if s.empty:  

                tail_means_k[name] = float('nan')  

                continue  

            # Tail mean (steady-state) for CV table & S6 export.  

            seg = s.iloc[cutoff:] if cutoff < len(s) else s.iloc[max(0, len(s)-1):]  

            tail_means_k[name] = float(seg.mean()) if len(seg) else float('nan')  

            # Smoothed/decimated trace for fan chart - default-config runs only.  

            if not entry['has_overrides']:  

                sm = s.rolling(20, min_periods=1).mean().values  

                sm_dec, _ = decimate_series(sm, MAX_PLOT_POINTS)  

                headline_traces_default[name].append(sm_dec)  

        headline_tail_means[k] = tail_means_k  

        # ---- S1 feature rows -----  

        ep_c = ep_df.iloc[cutoff:] if cutoff < n_ep else ep_df.iloc[max(0, n_ep-1):]  

        yr_c = yr_df[yr_df['episode'] >= cutoff] if (not yr_df.empty and 'episode'  

        ↪in yr_df.columns) else yr_df  

        cfg = entry['config']  

        for i in range(N_AGENTS):  

            row = {'sweep': k[0], 'variant': k[1], 'seed': k[2],  

                'agent': i + 1,  

                'archetype': archetype_of(i, cfg),  

                'weighting': weighting_of(i, cfg),  

                'has_overrides': entry['has_overrides']}  

            for prefix in EPISODE_FEATURES:  

                col = f'{prefix}_A{i+1}'  

                row[prefix] = float(ep_c[col].mean()) if col in ep_c.columns else  

        ↪np.nan  

            for prefix in YEAR_FEATURES:  

                col = f'{prefix}_A{i+1}'  

                row[prefix] = float(yr_c[col].mean()) if (not yr_c.empty and col in  

        ↪yr_c.columns) else np.nan  

            udbc_total = 0.0; udbc_vec = {}  

            for b in UDBC_BUCKETS:  

                col = f'udbc_{b}_total_A{i+1}'  

                v = float(ep_c[col].mean()) if col in ep_c.columns else 0.0  

                udbc_vec[b] = v; udbc_total += v  

            for b in UDBC_BUCKETS:

```

```

        row[f'udbc_{b}_share'] = (udbc_vec[b] / udbc_total) if udbc_total > 1e-9 else 0.0
        feat_rows.append(row)

# ---- SS2 outcome compare (green/penalty/clearing tail means) ----
green_means_per_run[k] = {
    'green_frac': tail_means_k.get('mean_green_frac', np.nan),
    'penalty': tail_means_k.get('mean_penalty', np.nan),
    'clearing_price': float(ep_c['clearing_price_last'].mean())
                        if 'clearing_price_last' in ep_c.columns and not ep_c.
empty else np.nan,
}

# ---- SS3.a anchor (per-(run, year) clearing in tail) ----
if not yr_c.empty and 'year' in yr_c.columns and 'clearing_price' in yr_c.
columns:
    cp_by_year = yr_c.groupby('year')['clearing_price'].mean()
    for y in range(N_YEARS):
        clearing_per_year_rows.append({
            'sweep': k[0], 'variant': k[1], 'seed': k[2],
            'has_overrides': entry['has_overrides'],
            'year': y,
            'clearing': float(cp_by_year.get(y, np.nan)),
            'anchor': compute_fundamental_anchor(y, cfg),
        })

# ---- SS3 H1/H2 panel (per (run, agent, year)) ----
if not yr_c.empty and 'year' in yr_c.columns:
    ymean = yr_c.groupby('year').mean(numeric_only=True)
    for i in range(N_AGENTS):
        inv_col = f'invest_cost_A{i+1}'
        bank_col = f'bank_start_A{i+1}'
        if inv_col not in ymean.columns:
            continue
        mix = cfg['companies']['initial_mix'][i]
        df_i = pd.DataFrame({
            'year': ymean.index.astype(int),
            'invest_cost': ymean[inv_col].values,
            'bank_start': ymean[bank_col].values if bank_col in ymean.
columns else np.nan,
            'clearing_price': ymean['clearing_price'].values
                        if 'clearing_price' in ymean.columns else np.
nan,
        })
        df_i['lag_clearing_price'] = df_i['clearing_price'].shift(1)
        df_i['sweep'] = k[0]; df_i['variant'] = k[1]; df_i['seed'] = k[2]
        df_i['agent'] = i + 1

```

```

        df_i['archetype'] = archetype_of(i, cfg)
        df_i['weighting'] = weighting_of(i, cfg)
        for j, t in enumerate(['coal', 'gas', 'onshore', 'offshore', '
↳solar']):
            df_i[f'init_{t}'] = float(mix[j])
            pay_pieces.append(df_i)

# ---- $R2 UDBC labels -----
if not yr_c.empty and 'year' in yr_c.columns:
    udbc_pieces.append(_udbc_panel_for_run(yr_c, cfg, k))

# ---- $R1.b outcome row -----
out_row = {'sweep': k[0], 'variant': k[1], 'seed': k[2]}
out_row['mean_clearing_price'] = (float(yr_c['clearing_price'].mean())
↳if (not yr_c.empty and 'clearing_price'
in yr_c.columns)
                                else np.nan)
last50 = ep_df.iloc[-50:] if n_ep else ep_df
out_row['mean_green_frac'] = float(np.nanmean([
    last50[f'green_frac_A{i+1}'].mean()
    for i in range(N_AGENTS) if f'green_frac_A{i+1}' in last50.columns
])) if not last50.empty else np.nan
out_row['total_penalty'] = float(np.nansum([
    last50[f'penalty_A{i+1}'].mean()
    for i in range(N_AGENTS) if f'penalty_A{i+1}' in last50.columns
])) if not last50.empty else np.nan
out_row['quality_score'] = (float(last50['quality_score'].mean())
↳if 'quality_score' in last50.columns else np.
nan)
out_rows.append(out_row)

# ---- $6 export row -----
export_rows.append({
    'sweep': k[0], 'variant': k[1], 'seed': k[2],
    'n_episodes': n_ep,
    'convergence_cutoff': int(cutoff),
    'tail_frac': round(1.0 - cutoff / max(1, n_ep), 3),
    'system_reward_mean': tail_means_k.get('system_reward', np.nan),
    'quality_score_mean': tail_means_k.get('quality_score', np.nan),
    'clearing_price_last_mean': tail_means_k.get('clearing_price_last', np.
↳nan),
    'mean_green_frac': tail_means_k.get('mean_green_frac', np.nan),
    'mean_penalty': tail_means_k.get('mean_penalty', np.nan),
    'mean_invest_cost': tail_means_k.get('mean_invest_cost', np.nan),
})

# Drop everything before the next iteration. The generator pattern

```

```

# would do this implicitly; we're explicit here to make the memory
# contract obvious.
del ep_df, yr_df, ep_c, yr_c, rew, px
gc.collect()

# Materialise the row-list artefacts into DataFrames once.
conv_df          = pd.DataFrame(conv_results)
feat_df          = pd.DataFrame(feat_rows)
clearing_per_year_df = pd.DataFrame(clearing_per_year_rows)
pay_df           = pd.concat(pay_pieces, ignore_index=True) if pay_pieces
↳ else pd.DataFrame()
udbc_df          = pd.concat(udbc_pieces, ignore_index=True) if udbc_pieces
↳ else pd.DataFrame()
out_df           = pd.DataFrame(out_rows)
del feat_rows, clearing_per_year_rows, pay_pieces, udbc_pieces, out_rows
gc.collect()

print(f'Pre-compute finished in {time.time() - _t0:.1f}s.')
print(f' feat_df: {feat_df.shape}      clearing_per_year_df:↳
↳ {clearing_per_year_df.shape}')
print(f' pay_df: {pay_df.shape}      udbc_df: {udbc_df.shape}      out_df:↳
↳ {out_df.shape}')

```

Streaming pre-compute over 20 run(s) (tail_episodes=None, peak RAM = 1 run)...

pre-compute: 0%| | 0/20 [00:00<?, ?run/s]

Pre-compute finished in 301.7s.

feat_df: (160, 30)	clearing_per_year_df: (240, 7)
pay_df: (1920, 16)	udbc_df: (60765696, 11) out_df: (20, 7)

1.7 §C0 — Sizing & Sanity Check

A quick readout of the loaded dataset and a coarse runtime estimate. Any section that scales with episode count is decimated to `MAX_DETECTOR_POINTS` (detection) or `MAX_PLOT_POINTS` (plotting) — those constants are defined in §5 and can be tuned globally.

```

[8]: # Dataset sizing + coarse ETA. Streaming pre-compute already ran above;
# everything here reads from cached scalars (no DataFrames touched).
_eps      = list(n_episodes_per_run.values()) or [0]
_max_ep   = max(_eps)
_total_ep_rows = sum(_eps)
_total_yr_rows = len(udbc_df) if not udbc_df.empty else 0 # post-tail estimate

print(f'Loaded {len(runs)} run(s) into the streaming registry.')
print(f' episodes/run: max={_max_ep:,} total ep rows: {_total_ep_rows:,}')

```

```

print(f'  decimation budgets: detector={MAX_DETECTOR_POINTS:,}  '
      f'  plot={MAX_PLOT_POINTS:,}')
print(f'  pre-compute tail cap: {PRECOMPUTE_TAIL_EPISODES}  '
      f'  (None = read full episode logs; integer = read last N episodes)')

# Heuristic ETA for the *downstream* sections - the heavy lifting is done.
_eta = {
    'C1 detector table': 1,
    'C1 plot':          max(2, len(runs) * 1),
    'C2 fan + CV table': max(2, len(runs) * 1),
    'S1 cluster + PCA': 3,
    'S3 H1/H2 fits':    3,
    'R1 cluster shifts': 3,
    'R2 logit + RF':     max(10, len(udbc_df) // 20_000),
}
print('\nCoarse per-section ETA (seconds, downstream only):')
for k, v in _eta.items():
    print(f'  {k:<22s} ~ {v}s')

```

Loaded 20 run(s) into the streaming registry.

episodes/run: max=120,000 total ep rows: 2,400,000

decimation budgets: detector=5,000 plot=4,000

pre-compute tail cap: None (None = read full episode logs; integer = read last N episodes)

Coarse per-section ETA (seconds, downstream only):

C1 detector table	~ 1s
C1 plot	~ 20s
C2 fan + CV table	~ 20s
S1 cluster + PCA	~ 3s
S3 H1/H2 fits	~ 3s
R1 cluster shifts	~ 3s
R2 logit + RF	~ 3038s

1.8 §C1 — Convergence (claim RQ1.a)

Two detectors on episode-level system reward and `clearing_price_last`: a curriculum-native rolling-mean plateau check, and the specialised PELT change-point. We use the **later** of the two as the conservative convergence cutoff per run, and flag detector disagreement.

The cutoff is then propagated to every steady-state metric in §C2–§S3.

```

[9]: # §C1 detector results - the rolling-plateau and PELT detectors ran
      # during the streaming pre-compute. Display the table here.
print(f'§C1 - convergence cutoffs over {len(conv_df)} run(s) '
      f'  (decimated to <={MAX_DETECTOR_POINTS} points each):')

```

```
display(conv_df)
```


SC1 - convergence cutoffs over 20 run(s) (decimated to <=5000 points each):

	run	n_ep	plateau_reward	plateau_price	pelt_reward	┐
└pelt_price	convergence_cutoff	tail_frac				
0	default/s1729	120000	NaN	NaN	78720	┐
└58440	78720	0.344				
1	default/s8191	120000	NaN	NaN	47160	┐
└78360	78360	0.347				
2	default/s6561	120000	NaN	NaN	78960	┐
└78240	78960	0.342				
3	default/s5041	120000	NaN	NaN	61320	┐
└55800	61320	0.489				
4	lrf_low/s1729	120000	NaN	NaN	93000	┐
└93960	93960	0.217				
5	lrf_low/s8191	120000	NaN	NaN	92160	┐
└92160	92160	0.232				
6	lrf_low/s6561	120000	NaN	NaN	87720	┐
└88440	88440	0.263				
7	lrf_high/s1729	120000	NaN	72480.0	109680	┐
└57960	109680	0.086				
8	lrf_high/s8191	120000	NaN	24456.0	104760	┐
└18480	104760	0.127				
9	lrf_high/s6561	120000	NaN	101136.0	92760	┐
└99120	101136	0.157				
10	no_msr/s1729	120000	NaN	NaN	87000	┐
└88080	88080	0.266				
11	no_msr/s8191	120000	NaN	NaN	95160	┐
└94920	95160	0.207				
12	no_msr/s6561	120000	84168.0	65136.0	81720	┐
└60120	84168	0.299				
13	all_fin/s1729	120000	NaN	NaN	37200	┐
└99240	99240	0.173				
14	all_fin/s8191	120000	NaN	NaN	76560	┐
└36960	76560	0.362				
15	all_fin/s6561	120000	NaN	NaN	84600	┐
└69120	84600	0.295				
16	balanced/s1729	120000	NaN	NaN	83760	┐
└77880	83760	0.302				
17	balanced/s8191	120000	NaN	NaN	102240	┐
└44760	102240	0.148				
18	balanced/s6561	120000	NaN	NaN	75360	┐
└88800	88800	0.260				
19	high_inf/s1729	120000	NaN	NaN	76920	┐
└73200	76920	0.359				

```

[10]: # Visualise reward and price trajectories with the agreed cutoff marked.
# Reads decimated arrays from `plot_traces` (built in pre-compute) so this
# cell never touches raw episode logs.
n = len(runs)
fig, axes = plt.subplots(n, 2, figsize=(14, 3.0 * n), squeeze=False)
for row_idx, k in enumerate(runs.keys()):
    tr = plot_traces.get(k, {})
    cutoff = tr.get('cutoff', 0)
    for ax, name, ylabel in [(axes[row_idx, 0], 'reward', 'system reward'),
                             (axes[row_idx, 1], 'price', 'clearing price (last_
→yr)')]:
        sub = tr.get(name)
        if not sub:
            ax.set_axis_off(); continue
        ax.plot(sub['x_dec'], sub['y_dec'], lw=0.6, alpha=0.5, color='#888',
                label='raw (decimated)')
        ax.plot(sub['sx_dec'], sub['sm_dec'], lw=1.6, color='C0',
                label='rolling mean')
        ax.axvline(cutoff, color='red', ls='--', lw=1.2,
                label=f'cutoff @ ep {cutoff}')
        ax.set(xlabel='episode', ylabel=ylabel, title=f'{k[1]}/s{k[2]}')
        ax.legend(fontsize=8, loc='lower right')
plt.tight_layout(); plt.show()

```



Interpretation. The two detectors should *broadly* agree on where the trajectory levels off; we take the later of the two as the conservative cutoff. With a 100-episode smoke run the detectors are noisy by construction (insufficient samples for a stable rolling mean or PELT prior), so the cutoff lands near the run end — that is expected. On the production-scale sweep these cutoffs typically settle around 60–80 % of `n_episodes`.


1.9 §C2 — Cross-Seed Reproducibility (claim RQ1.b)

Seed-bin fan charts on six headline metrics + a converged-window CV table. If converged-window CV is small relative to expected cross-variant differences, RQ2 / RQ3 claims are well-supported.

```
[11]: # Cross-seed fan charts from cached, smoothed+decimated headline traces
# (default-config runs only - cross-variant fan charts belong in SR1).
def _fan_panel(ax, name, traces):
    if not traces:
        ax.set_axis_off(); return
    L = min(len(s) for s in traces)
    arr = np.stack([s[:L] for s in traces])
    x = np.arange(L)
    med = np.median(arr, axis=0)
    lo, hi = np.percentile(arr, [16, 84], axis=0)
    ax.fill_between(x, lo, hi, alpha=0.25, color='C0', label='+/-1 sigma across
↳seeds')
    ax.plot(x, med, color='C0', lw=1.6, label='median')
    for s in arr:
        ax.plot(x, s, color='C0', lw=0.4, alpha=0.4)
    ax.set(title=name, xlabel='decimated step')
    ax.legend(fontsize=8)

fig, axes = plt.subplots(2, 3, figsize=(18, 8))
for ax, (name, *_ ) in zip(axes.flat, HEADLINE_SPEC):
    _fan_panel(ax, name, headline_traces_default.get(name, []))
plt.suptitle('§C2 - Cross-seed fan charts (default-config runs)', y=1.02,
↳fontsize=13)
plt.tight_layout(); plt.show()
```

§C2 - Cross-seed fan charts (default-config runs)



```
[12]: # Converged-window CV table from cached tail means (default-config runs).
default_keys = list(DEFAULT_RUNS)
rows = []
for name, *_ in HEADLINE_SPEC:
    vals = [headline_tail_means[k][name] for k in default_keys
            if k in headline_tail_means and np.isfinite(headline_tail_means[k].
    ↪get(name, np.nan))]
    if not vals: continue
    arr = np.array(vals, dtype=float)
    mu = float(arr.mean()); sd = float(arr.std(ddof=1)) if len(arr) > 1 else
    ↪float('nan')
    rows.append({'metric': name, 'n_seeds': len(arr),
                'mean across seeds': mu, 'std across seeds': sd,
                'cv (%)': 100.0 * sd / abs(mu) if mu else float('nan')})
cv_df = pd.DataFrame(rows)
display(cv_df.style.format({'mean across seeds': '{:.4g}', 'std across seeds':
    ↪'{:.3g}', 'cv (%)': '{:.2f}'}))
```

<pandas.io.formats.style.Styler at 0x2006dfe6c30>

Interpretation. On the production sweep we want headline cv (%) below ~10 %. On the smoke dataset (only 2 seeds × 100 episodes) the CVs are large and uninterpretable — the table is shown as a sanity check that the pipeline runs, not as evidence.

1.10 §S1 — Strategy Archetypes in the Converged Market (claim RQ3.a)

Per-(run, agent) feature vector built from the converged window. We cluster with KMeans (k chosen by **Silhouette + Elbow**), confirm with hierarchical clustering (Ward linkage, dendrogram), visualise in 2D via PCA, and cross-tabulate the emergent clusters against the **designed** archetype ×

reward-weighting cells. The clustering pipeline produced here is reused by §S2 (variant comparison) and §R1 (regulatory shifts).

```
[13]: # Per-(run, agent) feature matrix `feat_df` was built during the
# streaming pre-compute. Listed here are the column families it carries.
EPISODE_FEATURES = ['bid_price', 'avg_bid_mult', 'inv_onshore_share',
                    'inv_offshore_share', 'inv_solar_share', 'green_frac',
                    'invest_cost', 'penalty',
                    'peak_loan_outstanding', 'peak_carry_forward',
                    'final_treasury_reserve']
YEAR_FEATURES     = ['sec_buy_vol', 'sec_sell_vol',
                    'coverage_gap', 'effective_penalty_rate',
                    'carry_forward_end', 'treasury_reserve',
                    'loan_outstanding']
UDBC_BUCKETS      = ['U', 'D', 'M', 'B', 'C']

def feature_columns(df):
    base = list(EPISODE_FEATURES) + list(YEAR_FEATURES) + [f'udbc_{b}_share'
    ↪for b in UDBC_BUCKETS]
    return [c for c in base if c in df.columns]

# Filter feat_df to default-config rows for §S1 cluster fit; the override
# rows (used by §S2 / §R1.a) are projected into the same basis below.
default_feat_df = feat_df[~feat_df['has_overrides']].reset_index(drop=True)
print(f'Per-(run, agent) feature matrix - default rows: {default_feat_df.
    ↪shape[0]} '
      f'rows x {default_feat_df.shape[1]} cols '
      f'(full feat_df incl. overrides: {feat_df.shape[0]} rows)')
display(default_feat_df.head())
```

Per-(run, agent) feature matrix - default rows: 32 rows x 30 cols (full feat_df incl. overrides: 160 rows)

	sweep	variant	seed	agent	archetype	weighting	
	↪has_overrides	bid_price	avg_bid_mult	inv_onshore_share	inv_offshore_share		
	↪inv_solar_share	green_frac	invest_cost	\			
0	default_seeds	default	1729	1	coal-leaning	pure_financial	
	↪False	153.203461	0.510162		0.007709	0.980412	
	↪0.011878	0.607022	1072.403076				
1	default_seeds	default	1729	2	coal-leaning	balanced_esg	
	↪False	149.098495	0.509982		0.460026	0.000591	
	↪0.539383	0.605803	844.790283				
2	default_seeds	default	1729	3	balanced	pure_financial	
	↪False	50.001007	1.979976		0.007711	0.980583	
	↪0.011705	0.578509	22.947720				

3	default_seeds	default	1729	4	balanced	balanced_esg	
	False	86.572556	1.973350		0.987345	0.000589	
	0.012066	0.767926	1546.850830				
4	default_seeds	default	1729	5	transitioner	pure_financial	
	False	155.474350	1.253331		0.007598	0.980511	
	0.011891	0.730323	27.707409				
	penalty	peak_loan_outstanding	peak_carry_forward				
	final_treasury_reserve	sec_buy_vol	sec_sell_vol	coverage_gap			
	effective_penalty_rate	carry_forward_end	treasury_reserve				
0	1561.986328	0.0	1.942177	998.			
	697388	NaN	NaN	1.865883	155.08432		
	0.825303	879.442810					
1	132.408112	0.0	0.467835	1141.			
	180664	NaN	NaN	1.475546	155.08432		
	0.065799	907.337830					
2	109.117081	0.0	0.423553	971.			
	141663	NaN	NaN	1.795612	155.08432		
	0.055599	840.236938					
3	69.016022	0.0	0.297686	1166.			
	110352	NaN	NaN	0.353330	155.08432		
	0.035683	923.249512					
4	56.970215	0.0	0.234430	986.			
	537415	NaN	NaN	0.252307	155.08432		
	0.028163	824.462463					
	loan_outstanding	udbc_U_share	udbc_D_share	udbc_M_share	udbc_B_share		
	udbc_C_share						
0	0.0	0.155147	0.123360	0.288565	0.216183		
	0.216746						
1	0.0	0.066429	0.002261	0.031111	0.518246		
	0.381952						
2	0.0	0.082597	0.001309	0.026785	0.258173		
	0.631136						
3	0.0	0.167538	0.016426	0.031439	0.623673		
	0.160923						
4	0.0	0.097402	0.010503	0.031509	0.849044		
	0.011542						

```
[14]: # Standardise + KMeans with Silhouette/Elbow on default-config rows only.
def _prepare_X(df):
    feats = feature_columns(df)
    X = df[feats].values.astype(float)
    keep = []
    for j in range(X.shape[1]):
        col = X[:, j]
```

```

        col = col[np.isfinite(col)]
        if len(col) > 0 and np.nanstd(col) > 1e-12:
            keep.append(j)
    feats = [feats[j] for j in keep]
    X = X[:, keep]
    col_means = np.nanmean(X, axis=0)
    nan_mask = ~np.isfinite(X)
    X[nan_mask] = np.take(col_means, np.where(nan_mask)[1])
    scaler = StandardScaler().fit(X)
    return scaler.transform(X), feats, scaler

X, feats, scaler = _prepare_X(default_feat_df)
print('Feature columns kept after variance filter:', feats)
print('X shape:', X.shape)


K_RANGE = range(2, min(8, len(X)))
sil_scores, inertias = [], []
for k in K_RANGE:
    if k >= len(X):
        sil_scores.append(np.nan); inertias.append(np.nan); continue
    km = KMeans(n_clusters=k, n_init=10, random_state=0).fit(X)
    inertias.append(km.inertia_)
    try:
        sil_scores.append(silhouette_score(X, km.labels_))
    except Exception:
        sil_scores.append(np.nan)

fig, axes = plt.subplots(1, 2, figsize=(12, 4))
axes[0].plot(list(K_RANGE), inertias, 'o-')
axes[0].set(title='Elbow - KMeans inertia', xlabel='k', ylabel='inertia')
axes[1].plot(list(K_RANGE), sil_scores, 'o-', color='C2')
axes[1].set(title='Silhouette - KMeans', xlabel='k', ylabel='silhouette')
plt.tight_layout(); plt.show()

valid = [(k, s) for k, s in zip(K_RANGE, sil_scores) if np.isfinite(s)]
K_BEST = max(valid, key=lambda kv: kv[1])[0] if valid else 3
print(f'Selected k = {K_BEST} (max silhouette).')
```

Feature columns kept after variance filter: ['bid_price', 'avg_bid_mult', 'inv_onshore_share', 'inv_offshore_share', 'inv_solar_share', 'green_frac', 'invest_cost', 'penalty', 'peak_loan_outstanding', 'peak_carry_forward', 'final_treasury_reserve', 'coverage_gap', 'effective_penalty_rate', 'carry_forward_end', 'treasury_reserve', 'loan_outstanding', 'udbc_U_share', 'udbc_D_share', 'udbc_M_share', 'udbc_B_share', 'udbc_C_share']

X shape: (32, 21)




Selected k = 2 (max silhouette).

```
[15]: # Final KMeans + hierarchical confirmation on default-config rows.
km = KMeans(n_clusters=K_BEST, n_init=20, random_state=0).fit(X)
default_feat_df = default_feat_df.copy()
default_feat_df['kmeans_cluster'] = km.labels_

Z = linkage(X, method='ward')
hier_labels = fcluster(Z, t=K_BEST, criterion='maxclust')
default_feat_df['hier_cluster'] = hier_labels

fig, axes = plt.subplots(1, 2, figsize=(15, 5))
labels = [f"{r['variant'][:6]}/{s[r['seed']]}A{r['agent']}" for _, r in
↳ default_feat_df.iterrows()]
dendrogram(Z, labels=labels, leaf_rotation=90, ax=axes[0])
axes[0].set(title=f'Hierarchical clustering (Ward) - cut at k={K_BEST}')
axes[0].tick_params(axis='x', labelsize=7)

pca = PCA(n_components=2, random_state=0).fit(X)
P = pca.transform(X)
ax = axes[1]
for c in range(K_BEST):
    sel = km.labels_ == c
    ax.scatter(P[sel, 0], P[sel, 1], label=f'cluster {c}', s=60, alpha=0.85)
ax.set(title=f'PCA projection - KMeans clusters (var explained: {pca.
↳ explained_variance_ratio_.sum():.2%})',
        xlabel=f'PC1 ({pca.explained_variance_ratio_[0]:.1%})',
        ylabel=f'PC2 ({pca.explained_variance_ratio_[1]:.1%})')
ax.legend()
plt.tight_layout(); plt.show()
```



```
[16]: ct = pd.crosstab(
    [default_feat_df['archetype'], default_feat_df['weighting']],
    default_feat_df['kmeans_cluster'],
    margins=True,
)
print('Designed archetype x weighting vs KMeans cluster:')
display(ct)

ct_h = pd.crosstab(
    [default_feat_df['archetype'], default_feat_df['weighting']],
    default_feat_df['hier_cluster'],
    margins=True,
)
print('Designed archetype x weighting vs Hierarchical cluster:')
display(ct_h)
```

Designed archetype x weighting vs KMeans cluster:

kmeans_cluster		0	1	All
archetype	weighting			
balanced	balanced_esg	4	0	4
	pure_financial	1	3	4
coal-leaning	balanced_esg	3	1	4
	pure_financial	0	4	4
green-leader	balanced_esg	4	0	4
	pure_financial	4	0	4
transitioner	balanced_esg	4	0	4
	pure_financial	4	0	4
All		24	8	32

Designed archetype x weighting vs Hierarchical cluster:

hier_cluster		1	2	All
archetype	weighting			
balanced	balanced_esg	4	0	4
	pure_financial	0	4	4

coal-leaning	balanced_esg	1	3	4
	pure_financial	0	4	4
green-leader	balanced_esg	4	0	4
	pure_financial	4	0	4
transitioner	balanced_esg	4	0	4
	pure_financial	4	0	4
All		21	11	32

Interpretation. A diagonal-ish cross-tab indicates the emergent clustering recovers the designed archetypes. Off-diagonal mass is itself a finding (e.g. archetype heterogeneity collapsing under a specific reward weighting, or two archetypes converging onto the same strategy). The hierarchical and KMeans tables should agree on the qualitative grouping — when they don’t, the strategy space is ambiguous and the headline claim needs hedging.

```
[17]: # Persist the fitted feature space so §S2 and §R1 can project new
# (variant, seed, agent) feature vectors into the *same* space.
S1_ARTIFACTS = dict(scaler=scaler, pca=pca, kmeans=km, feature_cols=feats,
    ↪k=K_BEST)
print('Stored S1 artefacts: scaler, PCA, KMeans, feature_cols.')
```

Stored S1 artefacts: scaler, PCA, KMeans, feature_cols.

1.11 §S4 — Strategy-Feature Correlation Heatmap


Per-(run, agent) feature correlations across the same matrix used in §S1. Useful for the thesis interpretation pass: tells us which strategy axes are near-collinear (so a single PC explains most of the relevant variance) and which are independent.

```
[18]: # Pearson correlation across §S1 features on default-config rows.
_corr_feats = feature_columns(default_feat_df)
if _corr_feats and len(default_feat_df) >= 4:
    C = default_feat_df[_corr_feats].corr()
    fig, ax = plt.subplots(figsize=(min(14, 0.45 * len(_corr_feats) + 4),
                                     min(12, 0.45 * len(_corr_feats) + 3)))
    sns.heatmap(C, annot=True, fmt='.2f', cmap='RdBu_r', vmin=-1, vmax=1,
                square=True, cbar_kws={'shrink': 0.7}, ax=ax,
                annot_kws={'size': 7})
    ax.set_title('§S4 - Pearson correlation of strategy features,
    ↪(per-(run,agent) cells)')
    plt.tight_layout(); plt.show()
    pairs = []
    for i, a in enumerate(_corr_feats):
        for j, b in enumerate(_corr_feats):
            if j <= i: continue
            r = float(C.iloc[i, j])
            if abs(r) >= 0.9:
```

```

        pairs.append({'feature_a': a, 'feature_b': b, 'pearson_r': round(r, 3)})
    if pairs:
        print('Near-collinear feature pairs (|r| >= 0.9):')
        display(pd.DataFrame(pairs).sort_values('pearson_r', key=abs, ascending=False))
    else:
        print('No feature pairs with |r| >= 0.9 - features are reasonably independent.')
else:
    print('$S4 skipped - feature matrix too small or empty.')

```



Near-collinear feature pairs ($|r| \geq 0.9$):

	feature_a	feature_b	pearson_r
1	penalty	carry_forward_end	1.000

3	peak_loan_outstanding	loan_outstanding	1.000
0	penalty	peak_carry_forward	0.956
4	peak_carry_forward	carry_forward_end	0.955
2	penalty	udbc_M_share	0.940
6	carry_forward_end	udbc_M_share	0.937
5	peak_carry_forward	udbc_M_share	0.935

1.12 §S2 — How Reward Weights Shape Strategy (claim RQ3.b)


This is the headline RQ3 analysis. We use **two complementary contrasts**:

1. **Within the default config:** every default run already mixes pure-financial agents (`reward_weights=[1,0]`) with balanced-ESG agents (`reward_weights=[0.5,0.5]`). We project them into the §S1 PCA space and compare centroid positions and within-group spread.
2. **Cross-variant (when available):** if the loaded sweeps include an `all_financial`-style variant (every learning agent with `[w_cost=1, w_green=0]`), we re-run the §S1 clustering on that variant's per-agent feature matrix and compare cluster count, spread in the §S1 PCA basis, and silhouette score against the default.

```
[19]: # Within-default contrast - pure_financial vs balanced_esg agents in PCA.
fig, axes = plt.subplots(1, 2, figsize=(14, 5))

ax = axes[0]
for w, color in [('pure_financial', 'C3'), ('balanced_esg', 'C2')]:
    sel = (default_feat_df['weighting'] == w).values
    if sel.sum():
        ax.scatter(P[sel, 0], P[sel, 1], label=w, color=color, s=70, alpha=0.85,
                   edgecolor='k', linewidth=0.4)
        cx, cy = P[sel, 0].mean(), P[sel, 1].mean()
        ax.scatter([cx], [cy], marker='X', s=260, color=color,
                   edgecolor='black',
                   linewidth=1.5, label=f'{w} centroid')
ax.set(title='Within-default - pure-financial vs balanced-ESG agents in §S1_
PCA',
       xlabel=f'PC1 ({pca.explained_variance_ratio_[0]:.1%})',
       ylabel=f'PC2 ({pca.explained_variance_ratio_[1]:.1%})')
ax.legend(fontsize=8)

ax = axes[1]
mu = default_feat_df.groupby('weighting')[feats].mean()
mu_std = (mu - mu.mean()) / (mu.std(ddof=0) + 1e-9)
mu_std.T.plot(kind='bar', ax=ax, color=['C2', 'C3'])
ax.axhline(0, color='k', lw=0.5)
ax.set(title='Standardised feature means by weighting',
       ylabel='z-score (across weighting groups)', xlabel='')
ax.legend(title='weighting')
plt.tight_layout(); plt.show()
```



```
[20]: # Cross-variant contrast. Both `feat_df` (full) and `green_means_per_run`
# (cached during pre-compute) are read from artefacts; no runs are reloaded.
def _matches(variant, *needles):
    return any(n in variant for n in needles)

af_keys = [k for k, v in OVERRIDE_RUNS.items()
            if _matches(v['variant'], 'all_fin', 'all_financial')]
bal_keys = [k for k, v in OVERRIDE_RUNS.items()
            if _matches(v['variant'], 'balanced')]

def _project_keys(key_list):
    """Slice feat_df to the given run keys, project into the $S1 basis."""
    if not key_list:
        return np.empty((0, 2)), np.empty((0, len(feats))), 0
    mask = feat_df.set_index(['sweep', 'variant', 'seed']).index.isin(key_list)
    sub = feat_df[mask]
    if sub.empty:
        return np.empty((0, 2)), np.empty((0, len(feats))), 0
    Xv = sub[feats].values.astype(float)
    col_means = np.nanmean(Xv, axis=0)
    nan_mask = ~np.isfinite(Xv)
    Xv[nan_mask] = np.take(col_means, np.where(nan_mask)[1])
    Xv_std = scaler.transform(Xv)
    return pca.transform(Xv_std), Xv_std, len(Xv)

def _silhouette_sweep(X_std, k_max=8):
    if len(X_std) < 4:
        return [], None, None
    rng = range(2, min(k_max, len(X_std)))
    sils = []
    for k in rng:
```

```

        if k >= len(X_std):
            sils.append(np.nan); continue
        km_ = KMeans(n_clusters=k, n_init=10, random_state=0).fit(X_std)
        try:
            sils.append(silhouette_score(X_std, km_.labels_))
        except Exception:
            sils.append(np.nan)
    valid = [(k, s) for k, s in zip(rng, sils) if np.isfinite(s)]
    if not valid:
        return sils, None, None
    k_best = max(valid, key=lambda kv: kv[1])[0]
    km_final = KMeans(n_clusters=k_best, n_init=20, random_state=0).fit(X_std)
    return sils, k_best, km_final

def _spread(P_):
    return float(np.mean(np.linalg.norm(P_ - P_.mean(axis=0), axis=1))) if len(P_) > 0 else float('nan')

summary_rows = [{'variant': 'default', 'n': len(P),
                  'silhouette': max(sil_scores) if sil_scores else np.nan,
                  'k_best': K_BEST, 'mean radius (PC1-PC2)': _spread(P)}]

panels = [('default', P, km.labels_, f'default - KMeans (k={K_BEST})')]

for label, key_list in [('all_financial', af_keys), ('balanced', bal_keys)]:
    if not key_list:
        print(f'$S2 - {label} variant not loaded - panel skipped.')
        continue
    Pv, Xv_std, nv = _project_keys(key_list)
    sils_v, k_v, km_v = _silhouette_sweep(Xv_std)
    panels.append((label, Pv,
                  km_v.labels_ if km_v is not None else np.zeros(nv, dtype=int),
                  f'{label} - KMeans (k={k_v})' if k_v else f'{label} (n={nv})'))
    summary_rows.append({
        'variant': label, 'n': nv,
        'silhouette': max(sils_v) if sils_v else np.nan,
        'k_best': k_v, 'mean radius (PC1-PC2)': _spread(Pv),
    })

n_panels = len(panels)
fig, axes = plt.subplots(1, n_panels, figsize=(5.0 * n_panels, 4.5),
                        sharex=True, sharey=True, squeeze=False)
for ax, (title_lab, points, lbls, title) in zip(axes[0], panels):

```

```

    if len(points) == 0:
        ax.set_axis_off(); continue
    for c in np.unique(lbls):
        sel = lbls == c
        ax.scatter(points[sel, 0], points[sel, 1], s=60, alpha=0.85,
        ↪label=f'cluster {c}')
        ax.scatter(points[:, 0].mean(), points[:, 1].mean(), marker='X', s=240,
                    color='k', edgecolor='white', linewidth=1.5, label=f'{title_lab}
        ↪centroid')
        ax.set(title=title, xlabel='PC1', ylabel='PC2')
        ax.legend(fontsize=7)
plt.suptitle('§§2 - Strategy space across reward-weight variants (§§1 PCA
        ↪basis)', y=1.02)
plt.tight_layout(); plt.show()


display(pd.DataFrame(summary_rows))
print('Headline RQ3.b read: lower silhouette / smaller PC-radius in
        ↪`all_financial`'
      ' (or `balanced` with all financial agents) versus `default` is evidence
        ↪that'
      ' the ESG component drives strategic heterogeneity.')

def _green_summary(key_list):
    rows = []
    for k in key_list:
        gm = green_means_per_run.get(k)
        if gm is None: continue
        rows.append({'variant': k[1], 'seed': k[2], **gm})
    return pd.DataFrame(rows)

outcomes = [_green_summary(list(DEFAULT_RUNS)).assign(group='default')]
if af_keys: outcomes.append(_green_summary(af_keys).
        ↪assign(group='all_financial'))
if bal_keys: outcomes.append(_green_summary(bal_keys).assign(group='balanced'))
if len(outcomes) > 1:
    out_cmp = (pd.concat(outcomes, ignore_index=True)
                .groupby('group')[['green_frac', 'penalty', 'clearing_price']]
                .agg(['mean', 'std', 'count']))
    print('Outcome comparison across reward-weighting variants:')
    display(out_cmp)

```

§S2 - Strategy space across reward-weight variants (§S1 PCA basis)



	variant	n	silhouette	k_best	mean radius (PC1-PC2)
0	default	32	0.324160	2	3.019395
1	all_financial	24	0.483594	2	2.739593
2	balanced	24	0.307683	2	5.008859

Headline RQ3.b read: lower silhouette / smaller PC-radius in `all\_financial` (or `balanced` with all financial agents) versus `default` is evidence that the ESG component drives strategic heterogeneity.

Outcome comparison across reward-weighting variants:

	green_frac			penalty			
	clearing_price						
	mean	std	count	mean	std	count	
↪ mean		std	count				
group							
all_financial	0.742215	0.041169	3	447.120870	195.119828	3	67.
↪383401	21.782186	3					
balanced	0.612144	0.031027	3	1404.938599	489.834600	3	93.
↪982388	58.831070	3					
default	0.756268	0.015576	4	403.625168	122.888930	4	69.
↪069240	17.013179	4					

Interpretation. The within-default contrast is the lower-bar evidence: even *inside* the same training run, balanced-ESG agents should sit further along whichever PC tracks green investment than their pure-financial twins. The cross-variant comparison is the stronger test: if the all-financial variant collapses to fewer clusters and a tighter PCA cloud, removing the ESG component is *causing* strategic homogenisation. If the spread does **not** collapse, agents found other reasons to differentiate (cost-structure heterogeneity, urgency scalars, etc.) — a finding worth reporting honestly.

1.13 §S3 — Are Strategies Economically Rational? (claim RQ3.c, Tier 2)

Three sub-analyses:

- **H1.** *Carbon price drives green investment.* Per-(run, agent, year): label = 1 if `invest_cost` > `threshold`. Features = lag-1 clearing price, archetype dummies, agent bank at start of


year, year. Logistic regression with **L1/LASSO** (Lecture 4 + 8) and a **Random Forest** for the partial-dependence cross-check on the lagged carbon price.

- **H2.** *High-fossil agents transition faster.* Label = 1 if the agent's `delta_green` over the episode exceeds the median. Features = initial mix, archetype dummies, mean clearing price seen by the agent. Same LASSO + RF setup.
- **Anchor check.** Realised mean clearing price per year (across runs) vs `compute_fundamental_anchor(year, config)`. One overlay plot, Pearson correlation, and mean absolute deviation.

```
[21]: # Anchor-vs-realised credibility plot from cached `clearing_per_year_df`.
anchor_df = clearing_per_year_df[~clearing_per_year_df['has_overrides']].copy()
panel = anchor_df.groupby('year').agg(
    realised_mean=('clearing', 'mean'),
    realised_std=('clearing', 'std'),
    anchor=('anchor', 'mean'),
).reset_index()

fig, ax = plt.subplots(figsize=(10, 5))
ax.plot(panel['year'], panel['anchor'],          '-o', color='black', lw=2,
        label='fundamental anchor')
ax.plot(panel['year'], panel['realised_mean'], '-s', color='C0', lw=2,
        label='realised mean (across seeds)')
ax.fill_between(panel['year'],
                panel['realised_mean'] - panel['realised_std'].fillna(0),
                panel['realised_mean'] + panel['realised_std'].fillna(0),
                alpha=0.20, color='C0', label='+/-1 sigma across seeds')
ax.set(xlabel='year', ylabel='EUR / tCO2',
       title='SS3 anchor check - realised clearing vs fundamental anchor')
ax.legend()
plt.tight_layout(); plt.show()

valid = panel.dropna(subset=['realised_mean', 'anchor'])
if len(valid) > 1:
    corr = float(np.corrcoef(valid['realised_mean'], valid['anchor'])[0, 1])
    mad = float(np.mean(np.abs(valid['realised_mean'] - valid['anchor'])))
    print(f'Pearson correlation (realised, anchor): {corr:+.3f}')
    print(f'Mean absolute deviation: {mad:.2f} EUR/t')
else:
    print('Not enough year-level data points to compute correlation.')
```



Pearson correlation (realised, anchor): +0.994
Mean absolute deviation: 34.21 EUR/t

```
[22]: # Per-(run, agent, year) panel for §S3 hypothesis tests was built during
# the streaming pre-compute and lives in `pay_df`. Default-config slice only
# (override variants enter SR1, not the credibility section).
pay = pay_df[pay_df.set_index(['sweep', 'variant', 'seed']).index.
↳isin(list(DEFAULT_RUNS))].copy()
print(f'Per-(agent, year) panel (default-config): {len(pay)} rows')
display(pay.head())
```

Per-(agent, year) panel (default-config): 384 rows

	year	invest_cost	bank_start	clearing_price	lag_clearing_price		
↳sweep	variant	seed	agent	archetype	weighting	init_coal	init_gas
↳	init_onshore	init_offshore	init_solar				
0	0	91.545982	1.846466	45.968960		NaN	
↳default_seeds	default	1729	1	coal-leaning	pure_financial		0.25
↳	0.3	0.2	0.15	0.1			
1	1	110.796150	0.601051	45.359337	45.968960		
↳default_seeds	default	1729	1	coal-leaning	pure_financial		0.25
↳	0.3	0.2	0.15	0.1			
2	2	91.193016	0.967672	49.629295	45.359337		
↳default_seeds	default	1729	1	coal-leaning	pure_financial		0.25
↳	0.3	0.2	0.15	0.1			
3	3	63.510624	0.635490	56.886814	49.629295		
↳default_seeds	default	1729	1	coal-leaning	pure_financial		0.25
↳	0.3	0.2	0.15	0.1			

4	4	73.214508	0.631276	64.336433	56.886814	0.25
↪	default_seeds	default	1729	1	coal-leaning	pure_financial
↪	0.3	0.2	0.15	0.1		

```
[23]: # --- (c) H1 - does carbon price drive green investment? -----
INV_THRESH = float(np.nanpercentile(pay['invest_cost'], 75)) if len(pay) else 0.
↪0
print(f'Invest threshold (top-quartile invest_cost): {INV_THRESH:.2f} M€')

H1 = pay.dropna(subset=['lag_clearing_price', 'bank_start']).copy()
H1['invests'] = (H1['invest_cost'] > INV_THRESH).astype(int)
arche_dummies = pd.get_dummies(H1['archetype'], prefix='arch', drop_first=True)
H1_X = pd.concat([H1[['lag_clearing_price', 'bank_start', 'year']],
↪arche_dummies], axis=1)
H1_y = H1['invests']

if H1_y.nunique() < 2 or len(H1_y) < 30:
    print('H1 skipped - insufficient label variation or rows.')
else:
    pipe = Pipeline([
        ('scale', StandardScaler()),
        ('lr', LogisticRegression(penalty='l1', solver='saga',
↪max_iter=2000, C=0.5, random_state=0)),
    ])
    cv_score = cross_val_score(pipe, H1_X, H1_y,
↪cv=StratifiedKFold(n_splits=min(5, H1_y.value_counts().min()), shuffle=True,
↪random_state=0), scoring='roc_auc')
    pipe.fit(H1_X, H1_y)
    coefs = pd.Series(pipe.named_steps['lr'].coef_[0], index=H1_X.columns).
↪sort_values(key=abs, ascending=False)
    print(f'H1 logistic LASSO - CV AUC = {cv_score.mean():.3f} ± {cv_score.
↪std():.3f}')
    print('Coefficients (standardised features):')
    display(coefs.to_frame('coef'))

    # Random Forest + partial dependence on lagged clearing price.
    rf = RandomForestClassifier(n_estimators=300, max_depth=None,
↪random_state=0, n_jobs=-1).fit(H1_X, H1_y)
    fi = pd.Series(rf.feature_importances_, index=H1_X.columns).
↪sort_values(ascending=False)
    print('Random forest feature importances:')
    display(fi.to_frame('importance'))

    pd_res = partial_dependence(rf, H1_X, ['lag_clearing_price'],
↪grid_resolution=40)
    fig, ax = plt.subplots(figsize=(8, 4))
    ax.plot(pd_res['grid_values'][0], pd_res['average'][0], lw=2)
```

```

ax.axvline(50, color='red', ls='--', label='theory threshold 50 €/t')
ax.set(xlabel='lag clearing price (EUR/t)', ylabel='P(invests | x)',
      title='H1 - RF partial dependence on lagged carbon price')
ax.legend(); plt.tight_layout(); plt.show()

```

Invest threshold (top-quartile invest_cost): 119.63 M€


H1 logistic LASSO - CV AUC = 0.729 ± 0.049

Coefficients (standardised features):

	coef
arch_green-leader	-1.036510
arch_coal-leaning	-0.319778
year	-0.301089
bank_start	0.297173
lag_clearing_price	-0.117193
arch_transitioner	-0.057624

Random forest feature importances:

	importance
bank_start	0.483094
lag_clearing_price	0.213681
year	0.132263
arch_green-leader	0.076138
arch_coal-leaning	0.056834
arch_transitioner	0.037990



```

[24]: # --- (d) H2 - heterogeneity: high-fossil agents transition faster ----
agent_panel = (
    pay.groupby(['variant', 'seed', 'agent', 'archetype', 'weighting'])

```

```

        .agg(mean_clearing=('clearing_price', 'mean'),
              total_invest=('invest_cost', 'sum'),
              init_coal=('init_coal', 'first'),
              init_gas=('init_gas', 'first'),
              init_onshore=('init_onshore', 'first'),
              init_offshore=('init_offshore', 'first'),
              init_solar=('init_solar', 'first'))
        .reset_index()
    )
    if len(agent_panel) > 0 and 'total_invest' in agent_panel.columns:
        med = float(agent_panel['total_invest'].median())
        agent_panel['transitions'] = (agent_panel['total_invest'] > med).astype(int)
        arche_d = pd.get_dummies(agent_panel['archetype'], prefix='arch',
        ↪drop_first=True)
        H2_X = pd.concat([agent_panel[['mean_clearing', 'init_coal', 'init_gas']],
        ↪arche_d], axis=1)
        H2_y = agent_panel['transitions']
        if H2_y.nunique() < 2 or len(H2_y) < 8:
            print('H2 skipped - insufficient sample size for cross-validated logit.
            ↪')
        else:
            pipe2 = Pipeline([('scale', StandardScaler()),
                              ('lr', LogisticRegression(penalty='l1',
            ↪solver='saga', max_iter=2000, C=0.5, random_state=0))])
            n_splits_h2 = max(2, min(5, int(H2_y.value_counts().min())))
            cv2 = cross_val_score(pipe2, H2_X, H2_y,
            ↪cv=StratifiedKFold(n_splits=n_splits_h2,
            ↪shuffle=True, random_state=0),
                              scoring='roc_auc')
            pipe2.fit(H2_X, H2_y)
            coefs2 = pd.Series(pipe2.named_steps['lr'].coef_[0], index=H2_X.
            ↪columns).sort_values(key=abs, ascending=False)
            print(f'H2 logistic LASSO - CV AUC = {cv2.mean():.3f} ± {cv2.std():.
            ↪3f}')
            print('Sign of init_coal positive  higher-fossil agents transition
            ↪faster (the theory prediction).')
            display(coefs2.to_frame('coef'))
    else:
        print('H2 skipped - agent panel empty.')

```

H2 logistic LASSO - CV AUC = 0.408 ± 0.130

Sign of init_coal positive higher-fossil agents transition faster (the theory prediction).

	coef
init_coal	0.088168
mean_clearing	0.000000

```

init_gas          0.000000
arch_coal-leaning 0.000000
arch_green-leader 0.000000
arch_transitioner 0.000000

```

Interpretation. A positive lag-clearing-price coefficient in H1 plus a step-up in the RF partial dependence near the theoretical 50 €/t LCOE crossover supports the *carbon price drives green investment* claim. A positive `init_coal` coefficient in H2 supports the *high-fossil agents transition faster* claim. The anchor overlay provides the credibility floor: if realised prices track the fundamental, the simulation is producing economically-grounded prices, which feeds RQ1 too.

1.14 §R1 — Regulatory Shifts (claim RQ2.a)

If the loaded sweeps include non-default variants (LRF / MSR / cap changes) we run §S1's clustering on each variant separately and project everything into the **default-fitted** §S1 PCA basis to make the visual comparison direct. Outcomes are summarised with cross-seed boxplots and IQR-based outlier flagging.

```

[25]: # Filter feat_df by (sweep, variant) and project into the §S1 basis.
variant_keys = sorted({(k[0], k[1]) for k in runs})
variant_keys_overrides = [(s, v) for (s, v) in variant_keys
                           if any(runs[k]['has_overrides'] for k in runs
                                   if k[0] == s and k[1] == v)]

if not variant_keys_overrides:
    print('No override variants loaded - §R1 cluster-shift comparison skipped.')
else:
    n_var = len(variant_keys)
    n_cols = min(n_var, 3); n_rows = int(np.ceil(n_var / n_cols))
    fig, axes = plt.subplots(n_rows, n_cols, figsize=(5.5 * n_cols, 5 *
↪n_rows), squeeze=False)
    for ax_idx, (sw, vn) in enumerate(variant_keys):
        ax = axes[ax_idx // n_cols, ax_idx % n_cols]
        sub = feat_df[(feat_df['sweep'] == sw) & (feat_df['variant'] == vn)]
        if sub.empty:
            ax.set_axis_off(); continue
        v_X = sub[feats].values.astype(float)
        col_means = np.nanmean(v_X, axis=0)
        nan_mask = ~np.isfinite(v_X)
        v_X[nan_mask] = np.take(col_means, np.where(nan_mask)[1])
        v_X_std = scaler.transform(v_X)
        v_P = pca.transform(v_X_std)
        labels_pred = km.predict(v_X_std)
        for c in np.unique(labels_pred):
            sel = labels_pred == c
            ax.scatter(v_P[sel, 0], v_P[sel, 1], s=60, alpha=0.85,
                      label=f'default-cluster {c}')


```

```

ax.scatter(v_P[:, 0].mean(), v_P[:, 1].mean(), marker='X', s=260,
color='k',
          edgecolor='white', linewidth=1.5, label='variant centroid')
ax.set(title=f'{vn}', xlabel='PC1', ylabel='PC2')
ax.legend(fontsize=7)
for ax_idx in range(len(variant_keys), n_rows * n_cols):
    axes[ax_idx // n_cols, ax_idx % n_cols].set_axis_off()
plt.suptitle('§R1 - Strategy positions in §S1 PCA basis, by variant', y=1.
02, fontsize=13)
plt.tight_layout(); plt.show()

```

§R1 - Strategy positions in §S1 PCA basis, by variant



```

[26]: # Outcome boxplots from cached `out_df` (built in pre-compute).
OUT_METRICS = ['mean_clearing_price', 'mean_green_frac', 'total_penalty',
quality_score']


```

```

OUT_METRICS = [m for m in OUT_METRICS if m in out_df.columns]

fig, axes = plt.subplots(1, len(OUT_METRICS), figsize=(5 * len(OUT_METRICS), 4))
if len(OUT_METRICS) == 1:
    axes = [axes]
for ax, name in zip(axes, OUT_METRICS):
    sns.boxplot(data=out_df, x='variant', y=name, ax=ax, showfliers=False)
    sns.stripplot(data=out_df, x='variant', y=name, ax=ax, color='black',
        size=4, alpha=0.6)
    for v in out_df['variant'].unique():
        sub = out_df[out_df['variant'] == v][name].dropna()
        if len(sub) >= 4:
            q1, q3 = sub.quantile([0.25, 0.75]); iqr = q3 - q1
            mask = (sub < q1 - 1.5 * iqr) | (sub > q3 + 1.5 * iqr)
            for v_idx, value in zip(sub.index[mask], sub[mask]):
                ax.annotate('!', xy=(list(out_df['variant'].unique()).index(v),
                    value),
                        color='red', fontsize=12, ha='center')
    ax.set_title(name); ax.tick_params(axis='x', rotation=20)
plt.suptitle('SR1 - Outcome distribution by variant (red ! = IQR outlier
    seed)', y=1.03)
plt.tight_layout(); plt.show()
display(out_df.drop(columns=['sweep', 'seed']).groupby('variant').agg(['mean',
    'std', 'count']))

```



mean_clearing_price				mean_green_frac				
↪total_penalty				quality_score				
	mean	std	count	mean	std	count		
↪	mean	std	count	mean	std	count		
variant								
all_fin	54.300434	9.184643	3	0.741434	0.042773	3		
↪3601.810303	1888.703418	3	1.250820	0.277148	3			
balanced	63.052872	22.620674	3	0.609126	0.031296	3		
↪11267.906250	3640.908113	3	0.191500	0.574242	3			
default	55.342065	10.650117	4	0.753379	0.016073	4		
↪3310.676819	1202.726576	4	1.326705	0.275635	4			

high_inf	80.698227	NaN	1	0.772743	NaN	1	␣
↪3797.364014	NaN	1	1.350620	NaN	1		
lrf_high	176.593775	29.278658	3	0.733499	0.013589	3	␣
↪8133.390299	3325.457176	3	-0.794253	0.552371	3		
lrf_low	49.657721	3.712193	3	0.758230	0.017635	3	␣
↪1240.744344	517.737453	3	1.661887	0.081775	3		
no_msr	49.476554	5.001632	3	0.750652	0.026639	3	␣
↪2282.278259	2191.733399	3	1.541780	0.199128	3		

1.14.1 §R1.b — Cross-variant outcome statistics

For each loaded override variant we compute the per-seed mean of the headline outcomes, the percentage change vs. the default-config mean, and Cohen’s d against the default seeds. With ~3 seeds per variant these are descriptive (no formal hypothesis test) but they make the regulatory channel explicit: which outcome moves the most under each variant?

```
[27]: # Cross-variant outcome statistics. Rebuilt from `out_df` (cell §R1) so we
# stay consistent with the boxplot panel above.
if 'out_df' in dir() and not out_df.empty:
    metrics_present = [m for m in ['mean_clearing_price', 'mean_green_frac',
                                  'total_penalty', 'quality_score']
                      if m in out_df.columns]
    default_mask = out_df['variant'] == 'default'
    if not default_mask.any():
        print('§R1.b - `default` variant not loaded; skipping cross-variant_
↪stats.')
    else:
        ref = out_df[default_mask][metrics_present]
        ref_mu = ref.mean(); ref_sd = ref.std(ddof=1)

        rows = []
        for v in sorted(out_df['variant'].unique()):
            sub = out_df[out_df['variant'] == v][metrics_present]
            if sub.empty:
                continue
            for m in metrics_present:
                v_mu = float(sub[m].mean())
                v_sd = float(sub[m].std(ddof=1)) if len(sub) > 1 else_
↪float('nan')
                d_mu = float(ref_mu[m])
                # Pooled SD for Cohen's d (treat unequal sample sizes safely).
                pooled = float(np.sqrt(((len(sub) - 1) * (v_sd or 0) ** 2 +
                                  (len(ref) - 1) * (ref_sd[m] or 0) ** 2))_
↪/
                                  max(1, len(sub) + len(ref) - 2))) if_
↪len(ref) + len(sub) > 2 else float('nan')
```

```

        d_eff = (v_mu - d_mu) / pooled if (pooled and np.
↳isfinite(pooled) and pooled > 1e-12) else float('nan')
        pct = 100.0 * (v_mu - d_mu) / d_mu if abs(d_mu) > 1e-9 else
↳float('nan')
        rows.append({'variant': v, 'metric': m, 'n_seeds':
↳int(len(sub)),
                    'mean': v_mu, 'std': v_sd, 'Δ vs default (%)': pct,
                    "Cohen's d": d_eff})
    rq2_stats = pd.DataFrame(rows)
    print('§R1.b - Per-variant outcome stats (vs. default reference):')
    display(rq2_stats.style.format({'mean': '{:.4g}', 'std': '{:.3g}',
        'Δ vs default (%)': '{:+.1f}',
        "Cohen's d": '{:+.2f}'}))

    # Identify the dominant regulatory channel per non-default variant.
    non_def = rq2_stats[rq2_stats['variant'] != 'default'].
↳dropna(subset=["Cohen's d"])
    if not non_def.empty:
        dominant = (non_def.assign(_abs_d=non_def["Cohen's d"].abs())
                    .sort_values('_abs_d', ascending=False)
                    .groupby('variant')
                    .head(1)
                    .drop(columns='_abs_d'))
        print('§R1.b - Largest-effect channel per variant '
              '(largest |Cohen\'s d| against default):')
        display(dominant.style.format({'mean': '{:.4g}', 'std': '{:.3g}',
            'Δ vs default (%)': '{:+.1f}',
            "Cohen's d": '{:+.2f}'}))
    else:
        print('§R1.b skipped - outcome frame `out_df` not built.')

```

§R1.b - Per-variant outcome stats (vs. default reference):

<pandas.io.formats.style.Styler at 0x201e6222270>

§R1.b - Largest-effect channel per variant (largest |Cohen's d| against default):

<pandas.io.formats.style.Styler at 0x201e89c52b0>

1.14.2 §R1.c — Inflation-stress anchor check (high_inf variant)

The thesis sweep includes a single-seed `high_inf` cell with elevated penalty inflation drift (4 % vs. 2 % default) and $\sim 1.5\times$ the inflation variance. Two design components are sensitive to inflation:

- The **inflation-aware MA3 budget gate** uses $\text{MA3} \times \text{infl}(t)/\text{infl}(t-1)$ so the price floor adapts to drift; a runaway clearing price under high inflation would mean the gate is failing.

- The **real-anchor ESG flow term** uses $\text{anchor_real_t} / \text{anchor_real_0}$, i.e. drift is divided out by construction.

We re-run §S3's anchor-credibility check restricted to **high_inf** and report the same correlation / MAD numbers as a stress sanity check.

```
[28]: # high_inf inflation-stress anchor check from cached clearing_per_year_df.
inf_keys = [k for k, v in runs.items() if 'high_inf' in v['variant']]
if not inf_keys:
    print('SR1.c - `high_inf` variant not loaded; inflation-stress check_
↳skipped.')
else:
    inf_df = clearing_per_year_df[
        clearing_per_year_df.set_index(['sweep', 'variant', 'seed']).index.
↳isin(inf_keys)
    ].copy()
    if inf_df.empty or inf_df['clearing'].dropna().empty:
        print('SR1.c - high_inf rows lack clearing-price data - skipping.')
    else:
        panel_inf = (inf_df.groupby('year')
                      .agg(realised_mean=('clearing', 'mean'),
                          realised_std=('clearing', 'std'),
                          anchor=('anchor', 'mean'))
                      .reset_index())
        fig, ax = plt.subplots(figsize=(10, 4.5))
        ax.plot(panel_inf['year'], panel_inf['anchor'],          '-o',
↳color='black', lw=2, label='fundamental anchor')
        ax.plot(panel_inf['year'], panel_inf['realised_mean'], '-s',
↳color='C3', lw=2, label='realised mean (high_inf seeds)')
        ax.fill_between(panel_inf['year'],
                        panel_inf['realised_mean'] - panel_inf['realised_std'].
↳fillna(0),
                        panel_inf['realised_mean'] + panel_inf['realised_std'].
↳fillna(0),
                        alpha=0.20, color='C3', label='+/-1 sigma across seeds')
        ax.set(xlabel='year', ylabel='EUR / tCO2',
                title='SR1.c - high_inf: realised clearing vs fundamental_
↳anchor')
        ax.legend(); plt.tight_layout(); plt.show()

        valid = panel_inf.dropna(subset=['realised_mean', 'anchor'])
        if len(valid) > 1:
            corr = float(np.corrcoef(valid['realised_mean'],
↳valid['anchor'])[0, 1])
            mad = float(np.mean(np.abs(valid['realised_mean'] -
↳valid['anchor'])))
```


```

print(f'high_inf - Pearson correlation (realised, anchor): {corr:+.
↪3f}')

print(f'high_inf - Mean absolute deviation: {mad:.2f} EUR/t')
print('Read: tracking should be comparable to the default-config
↪anchor check '

      'because the anchor is computed in real EUR by construction. '
      'A large gap indicates the inflation-aware budget gate is
↪binding.')

```



```

high_inf - Pearson correlation (realised, anchor): +0.979
high_inf - Mean absolute deviation: 22.10 EUR/t
Read: tracking should be comparable to the default-config anchor check because
the anchor is computed in real EUR by construction. A large gap indicates the
inflation-aware budget gate is binding.

```

1.15 §R2 — UDBC Compliance Pathways (claim RQ2.b)

The simulator already classifies each (agent, year) cell into one of five buckets — **U** (under-allocation), **D** (debt-cascade), **M** (mixed), **B** (bought-into-compliance), **C** (compliant). We exploit this directly:

- **Multinomial logistic regression** predicting the bucket from bank_start, cap, clearing_price, archetype, year, and variant. **Random Forest** as the non-linear robustness check.
- **Class imbalance:** C dominates → fit with and without **SMOTE** oversampling and report the comparison.
- **Transition heatmaps** of year-t → year-t+1 bucket transitions per variant.

Per-cell labels live in `udbc_*_total_A{i}` only as *episode totals*, so we proxy the per-(agent, year) bucket by combining year-log shortfall and the episode-level bucket totals: a year is B if the agent

traded volume on the secondary > 0 and shortfall = 0, C if shortfall = 0 and no buying, U otherwise (alloc < emiss, no inherited cf), D if there is positive carry-forward, M otherwise.

```
[29]: # UDBC panel was built during the streaming pre-compute.
print(f'UDBC panel: {len(udbc_df)} (run, agent, year) cells')
if not udbc_df.empty:
    print('Class distribution:')
    print(udbc_df['label'].value_counts())
```

UDBC panel: 60765696 (run, agent, year) cells

Class distribution:

label

C 37833010

B 11428138

U 6217831

M 5286717

Name: count, dtype: int64

label

C 37833010

B 11428138

U 6217831

M 5286717

Name: count, dtype: int64

```
[30]: # --- Multinomial logit with / without SMOTE -----
# On the production sweep this panel can hit tens of millions of rows. Cap
# the fit set with a stratified subsample (`R2_FIT_CAP`); the test set is
# also drawn from that stratified pool so the classification report stays
# representative of the converged-window distribution.
R2_FIT_CAP = 200_000

if udbc_df.empty or udbc_df['label'].nunique() < 2:
    print('UDBC modelling skipped - insufficient label variety.')
else:
    feat_cols = ['cap', 'clearing_price', 'bank_start', 'year']
    cat_cols = ['archetype', 'weighting', 'variant']
    df_m = udbc_df.dropna(subset=feat_cols + ['label']).copy()

    if len(df_m) > R2_FIT_CAP:
        # Stratified subsample preserves rare classes (esp. U/D/M).
        # Build per-class caps explicitly to avoid the pandas
        # DeprecationWarning around .apply on grouped frames.
        n_total = len(df_m)
        caps = {lab: max(50, R2_FIT_CAP * cnt // n_total)
                 for lab, cnt in df_m['label'].value_counts().items()}
        keep_idx = []
        for lab, sub in df_m.groupby('label'):
            n = min(len(sub), caps[lab])
```

```

        keep_idx.extend(sub.sample(n=n, random_state=0).index)
    df_m = df_m.loc[keep_idx].reset_index(drop=True)
    print(f'SR2 - subsampled fit set to {len(df_m)} rows (cap={R2_FIT_CAP}).
↳')

    Xm = pd.concat([df_m[feat_cols].astype(float),
                    pd.get_dummies(df_m[cat_cols], drop_first=True)], axis=1)
    ym = df_m['label'].values

    Xtr, Xte, ytr, yte = train_test_split(Xm, ym, test_size=0.25,
↳random_state=0, stratify=ym)
    pipe = Pipeline([('sc', StandardScaler(with_mean=False)),
                    ('lr', LogisticRegression(max_iter=2000, solver='lbfgs',
↳random_state=0, n_jobs=-1))])
    pipe.fit(Xtr, ytr)
    print('Multinomial logit (no SMOTE) - test classification report:')
    print(classification_report(yte, pipe.predict(Xte)))

    if _HAS_SMOTE and pd.Series(ytr).value_counts().min() >= 6:
        try:
            from imblearn.pipeline import Pipeline as ImbPipe
            sm_pipe = ImbPipe([('sc', StandardScaler(with_mean=False)),
                              ('sm', SMOTE(random_state=0, k_neighbors=3)),
                              ('lr', LogisticRegression(max_iter=2000,
↳solver='lbfgs', random_state=0, n_jobs=-1))])
            sm_pipe.fit(Xtr, ytr)
            print('Multinomial logit (with SMOTE) - test classification report:
↳')

            print(classification_report(yte, sm_pipe.predict(Xte)))
        except Exception as e:
            print(f'SMOTE step skipped: {e}')
        else:
            print('SMOTE skipped - not enough samples in the minority classes for
↳k-neighbours.')

    # Random Forest robustness + feature importances.
    rfm = RandomForestClassifier(n_estimators=400, random_state=0, n_jobs=-1).
↳fit(Xtr, ytr)
    print('Random forest - test classification report:')
    print(classification_report(yte, rfm.predict(Xte)))
    fi = pd.Series(rfm.feature_importances_, index=Xm.columns).
↳sort_values(ascending=False).head(10)
    display(fi.to_frame('importance'))

```

SR2 - subsampled fit set to 199997 rows (cap=200000).
Multinomial logit (no SMOTE) - test classification report:

precision	recall	f1-score	support
-----------	--------	----------	---------

B	0.00	0.00	0.00	9403
C	0.70	0.93	0.80	31131
M	0.70	0.09	0.16	4350
U	0.54	0.86	0.66	5116
accuracy			0.68	50000
macro avg	0.48	0.47	0.41	50000
weighted avg	0.55	0.68	0.58	50000

Multinomial logit (with SMOTE) - test classification report:

	precision	recall	f1-score	support
B	0.29	0.47	0.35	9403
C	0.79	0.45	0.57	31131
M	0.33	0.47	0.39	4350
U	0.48	0.99	0.64	5116
accuracy			0.51	50000
macro avg	0.47	0.59	0.49	50000
weighted avg	0.62	0.51	0.52	50000


Random forest - test classification report:

	precision	recall	f1-score	support
B	0.43	0.40	0.41	9403
C	0.76	0.77	0.76	31131
M	0.44	0.35	0.39	4350
U	0.60	0.72	0.66	5116
accuracy			0.66	50000
macro avg	0.56	0.56	0.56	50000
weighted avg	0.65	0.66	0.65	50000

	importance
bank_start	0.621788
clearing_price	0.155594
cap	0.035873
year	0.033394
weighting_pure_financial	0.032305
archetype_coal-leaning	0.020186
archetype_green-leader	0.018152
archetype_transitioner	0.015635
variant_balanced	0.014211
variant_no_msr	0.013321

```
[31]: # --- Transition heatmaps year-t → year-t+1 per variant -----
def _transition_matrix(df_var):
    pairs = []
    for (variant, seed, agent), grp in df_var.groupby(['variant', 'seed',
    ↪ 'agent']):
        grp_sorted = grp.sort_values('year')
        labs = grp_sorted['label'].values
        for a, b in zip(labs[:-1], labs[1:]):
            pairs.append((a, b))
    if not pairs:
        return pd.DataFrame(0, index=UDBC_BUCKETS, columns=UDBC_BUCKETS)
    df_p = pd.DataFrame(pairs, columns=['from', 'to'])
    M = pd.crosstab(df_p['from'], df_p['to'])
    M = M.reindex(index=UDBC_BUCKETS, columns=UDBC_BUCKETS, fill_value=0)
    return M.div(M.sum(axis=1).replace(0, np.nan), axis=0).fillna(0)

variants_present = sorted(udbc_df['variant'].unique())
n_v = len(variants_present)
fig, axes = plt.subplots(1, n_v, figsize=(4.2 * n_v, 4), squeeze=False)
for ax, v in zip(axes[0], variants_present):
    M = _transition_matrix(udbc_df[udbc_df['variant'] == v])
    sns.heatmap(M, annot=True, fmt='.2f', cmap='Blues', vmin=0, vmax=1, ax=ax,
    cbar=(ax is axes[0, -1]))
    ax.set(title=f'{v} - UDBC transitions (rows: from, cols: to)',
    xlabel='to', ylabel='from')
plt.suptitle('SR2 - UDBC year-t → year-t+1 transition probabilities', y=1.04)
plt.tight_layout(); plt.show()
```



Interpretation. Coefficients on `cap` (negative) and `clearing_price` (positive) drive the model toward U / D states — that’s the regulatory lever turning the screw. The transition heatmap diagonals tell us which buckets are *sticky*: tighter LRF should plausibly raise the U → U probability and lower the U → C recovery rate. Comparing variants side-by-side answers RQ2.b directly.

1.16 Summary — Mapping Panels Back to Claims

Claim	Section	Panels
RQ1.a — Agents converge by episode N	§C1	reward + price trajectory with PELT/plateau cutoff
RQ1.b — Findings replicate across seeds	§C2	fan charts + cross-seed CV table
RQ3.a — Distinguishable strategic archetypes	§S1	elbow/silhouette, dendrogram, PCA, archetype × cluster crosstab
RQ3.b — Removing ESG component collapses the strategy space	§S2	within-default contrast + all_fin & balanced PCA panels + spread/silhouette table + outcome comparison
RQ3.c — Agents respond rationally to incentives	§S3	anchor overlay, H1 LASSO + RF partial dependence, H2 logit
RQ2.a — Regulatory levers shift strategies and outcomes	§R1 + §R1.b	per-variant PCA panels + outcome boxplots + $\Delta\%$ /Cohen's d table
RQ2.a (inflation stress) — real-anchor & MA3 gate hold under elevated inflation	§R1.c	anchor vs. realised clearing on high_inf , with correlation / MAD
RQ2.b — Stricter regulation changes the <i>type</i> of compliance failure	§R2	multinomial logit (with/without SMOTE) + transition heatmaps

Reproducibility. Every section keys off (`sweep`, `variant`, `seed`) and writes nothing back to disk. Re-run from §1 to refresh.

Out of scope. Pathology / autoencoder analyses and the clearing-price prediction model dropped in v3 of the planning doc are deliberately not implemented here.

1.17 §Export — Headline Metrics Table

Persist a (sweep, variant, seed) → headline metrics table to `results/analysis/headline_metrics.csv`. As the thesis-experiments sweep adds more variants, this CSV becomes the convenient single-source-of-truth for cross-experiment comparison plots and tables.

```
[32]: # Cross-run headline export from cached `export_rows`. The streaming
# pre-compute already used the C1 convergence cutoff for tail means, so
# every value here is steady-state by construction.
_export_df = pd.DataFrame(export_rows).sort_values(['sweep', 'variant', 'seed'])
_out_dir = PROJECT_ROOT / 'results' / 'analysis'
_out_dir.mkdir(parents=True, exist_ok=True)
_out_path = _out_dir / 'headline_metrics.csv'
_export_df.to_csv(_out_path, index=False)
```

```
print(f'Wrote {len(_export_df)} rows -> {_out_path}')
display(_export_df)
```

Wrote 20 rows -> C:\Users\danie\Documents\Python_Scripts\Master Thesis\Thesis-Energy-Auction\results\analysis\headline_metrics.csv

	sweep	variant	seed	n_episodes	convergence_cutoff	
	tail_frac	system_reward_mean	quality_score_mean	clearing_price_last_mean		
	mean_green_frac	mean_penalty	mean_invest_cost			
0	default_seeds	default	1729	120000	78720	0.
	344	17.726650	1.702850	85.903183		
	0.737669	268.799408	681.168396			
3	default_seeds	default	5041	120000	61320	0.
	489	23.416452	1.026984	52.062107		
	0.749559	554.441467	657.331970			
2	default_seeds	default	6561	120000	78960	0.
	342	10.767038	1.367920	81.315308		
	0.771806	441.747559	856.779419			
1	default_seeds	default	8191	120000	78360	0.
	347	30.844290	1.287548	56.996361		
	0.766038	349.512238	812.467407			
13	thesis_experiments	all_fin	1729	120000	99240	0.
	173	18.047422	1.563283	65.162247		
	0.771296	272.863983	874.065002			
15	thesis_experiments	all_fin	6561	120000	84600	0.
	295	7.797767	1.051777	90.191063		
	0.695107	410.565643	247.831970			
14	thesis_experiments	all_fin	8191	120000	76560	0.
	362	33.245567	1.148012	46.796894		
	0.760243	657.932983	715.138794			
16	thesis_experiments	balanced	1729	120000	83760	0.
	302	-33.574993	0.378618	68.051613		
	0.617221	1184.304443	371.465118			
18	thesis_experiments	balanced	6561	120000	88800	0.
	260	-67.207787	-0.472465	161.324158		
	0.578892	1966.286011	25.823000			
17	thesis_experiments	balanced	8191	120000	102240	0.
	148	-20.346766	0.672953	52.571392		
	0.640320	1064.225342	535.516479			
19	thesis_experiments	high_inf	1729	120000	76920	0.
	359	3.081601	1.335410	163.130600		
	0.777294	422.724426	1056.660767			
7	thesis_experiments	lrf_high	1729	120000	109680	0.
	086	-103.951668	-0.879475	249.604065		
	0.739881	1092.279663	620.774109			

9	thesis_experiments	lrf_high	6561	120000	101136	0.
	↪157	-50.854996	-0.098338		240.210342	└
	↪0.722231	543.607483	439.474152			
8	thesis_experiments	lrf_high	8191	120000	104760	0.
	↪127	-99.228867	-1.128468		249.793198	└
	↪0.747509	1284.851440	642.496582			
4	thesis_experiments	lrf_low	1729	120000	93960	0.
	↪217	20.538864	1.662670		60.091877	└
	↪0.737664	129.197098	570.562195			
6	thesis_experiments	lrf_low	6561	120000	88440	0.
	↪263	28.686163	1.681295		46.698769	└
	↪0.773310	173.237106	842.648132			
5	thesis_experiments	lrf_low	8191	120000	92160	0.
	↪232	29.390514	1.798001		50.520416	└
	↪0.763020	91.008980	804.464050			
10	thesis_experiments	no_msr	1729	120000	88080	0.
	↪266	35.265640	1.505756		57.954685	└
	↪0.733182	209.534515	540.547058			
12	thesis_experiments	no_msr	6561	120000	84168	0.
	↪299	53.802654	1.766384		45.000900	└
	↪0.781162	21.351980	924.075378			
11	thesis_experiments	no_msr	8191	120000	95160	0.
	↪207	19.792583	1.497308		74.529976	└
	↪0.742212	531.967529	582.372803			